

MUltifrontal Massively Parallel Solver (MUMPS Version 4.7.1) Users' guide *

April 13, 2007

Abstract

This document describes the Fortran 90 and C user interface to MUMPS Version 4.7.1 We describe in detail the data structures, parameters, calling sequences, and error diagnostics. Example programs using MUMPS are also given.

*Information on how to obtain updated copies of MUMPS can be obtained from the Web pages <http://www.enseiht.fr/apo/MUMPS/> and <http://graal.ens-lyon.fr/MUMPS/> or by sending email to mumps@cerfacs.fr

Contents

1	Introduction	4
2	Main functionalities of MUMPS 4.7.1	5
2.1	Input matrix structure	5
2.2	Symmetric orderings	5
2.3	Other pre-processing facilities	5
2.4	Post-processing facilities	6
2.5	Solving the transposed system	6
2.6	Reduce/condense a problem on an interface (Schur complement and reduced/condensed right-hand side)	6
2.7	Arithmetic versions	7
2.8	The working host processor	8
2.9	Sequential version	8
2.10	Shared memory version	8
2.11	Main changes between versions	8
3	Sequence in which routines are called	8
4	Input and output parameters	11
4.1	Version number	11
4.2	Control of the three main phases: Analysis, Factorization, Solve	11
4.3	Control of parallelism	12
4.4	Matrix type	13
4.5	Centralized assembled matrix input: ICNTL(5)=0 and ICNTL(18)=0	13
4.6	Element matrix input: ICNTL(5)=1 and ICNTL(18)=0	14
4.7	Distributed assembled matrix input: ICNTL(5)=0 and ICNTL(18)≠0	14
4.8	Scaling	15
4.9	Given ordering: ICNTL(7)=1	15
4.10	Schur complement with reduced (or condensed) right-hand side: ICNTL(19) and ICNTL(26)	15
4.11	Workspace parameters	17
4.12	Right-hand side and solution vectors/matrices	18
5	Control parameters	19
6	Information parameters	26
6.1	Information local to each processor	26
6.2	Information available on all processors	27
7	Error diagnostics	28
8	Calling MUMPS from C	31
8.1	Array indices	32
8.2	Issues related to the C and Fortran communicators	32
8.3	Fortran I/O	32
8.4	Runtime libraries	33
8.5	Integer, real and complex datatypes in C and Fortran	33
8.6	Sequential version	33
9	Scilab and MATLAB interfaces	33
10	Examples of use of MUMPS	35
10.1	An assembled problem	35
10.2	An elemental problem	37
10.3	An example of calling MUMPS from C	37

1 Introduction

MUMPS (“MULTifrontal Massively Parallel Solver”) is a package for solving systems of linear equations of the form $\mathbf{Ax} = \mathbf{b}$, where $\mathbf{A}$ is a square sparse matrix that can be either unsymmetric, symmetric positive definite, or general symmetric. MUMPS uses a multifrontal technique which is a direct method based on either the LU or the LDL^T factorization of the matrix. We refer the reader to the papers [3, 4, 7, 16, 17, 21, 20, 8] for full details of the techniques used. MUMPS exploits both parallelism arising from sparsity in the matrix $\mathbf{A}$ and from dense factorizations kernels.

The main features of the MUMPS package include the solution of the transposed system, input of the matrix in assembled format (distributed or centralized) or elemental format, error analysis, iterative refinement, scaling of the original matrix, detection of zero pivots, and return of a Schur complement matrix. MUMPS offers several built-in ordering algorithms, a tight interface to some external ordering packages such as METIS [22] (strongly recommended) and PORD [23], and the possibility for the user to input a given ordering. Finally, MUMPS is available in various arithmetics (real or complex, single or double precision).

The software is written in Fortran 90 although a C interface is available (see Section 8). The parallel version of MUMPS requires MPI [24] for message passing and makes use of the BLAS [12, 13], BLACS, and ScaLAPACK [10] libraries. The sequential version only relies on BLAS.

MUMPS is downloaded from the web site almost once a day on average and has been run on very many machines, compilers and operating systems, although our experience is really only with UNIX-based systems. We have tested it extensively on parallel computers from SGI, Cray, and IBM and on clusters of workstations.

MUMPS distributes the work tasks among the processors, but an identified processor (the host) is required to perform most of the analysis phase, to distribute the incoming matrix to the other processors (slaves) in the case where the matrix is centralized, and to collect the solution. The system $\mathbf{Ax} = \mathbf{b}$ is solved in three main steps:

1. **Analysis.** The host performs an ordering (see Section 2.2) based on the symmetrized pattern $\mathbf{A} + \mathbf{A}^T$, and carries out symbolic factorization. A mapping of the multifrontal computational graph is then computed, and symbolic information is transferred from the host to the other processors. Using this information, the processors estimate the memory necessary for factorization and solution.
2. **Factorization.** The original matrix is first distributed to processors that will participate in the numerical factorization. The numerical factorization on each frontal matrix is conducted by a *master* processor (determined by the analysis phase) and one or more *slave* processors (determined dynamically). Each processor allocates an array for the so called contribution blocks and for the factors; the factors must be kept for the solution phase.
3. **Solution.** The right-hand side $\mathbf{b}$ is broadcast from the host to the other processors. These processors compute the solution $\mathbf{x}$ using the (distributed) factors computed during Step 2, and the solution is either assembled on the host or kept distributed on the processors.

Each of these phases can be called separately and several instances of MUMPS can be handled simultaneously. MUMPS allows the host processor to participate in computations during the factorization and solve phases, just like any other processor (see Section 2.8).

For both the symmetric and the unsymmetric algorithms used in the code, we have chosen a fully asynchronous approach with dynamic scheduling of the computational tasks. Asynchronous communication is used to enable overlapping between communication and computation. Dynamic scheduling was initially chosen to accommodate numerical pivoting in the factorization. The other important reason for this choice was that, with dynamic scheduling, the algorithm can adapt itself at execution time to remap work and data to more appropriate processors. In fact, we combine the main features of static and dynamic approaches; we use the estimation obtained during the analysis to map some of the main computational tasks; the other tasks are dynamically scheduled at execution time. The main data structures (the original matrix and the factors) are similarly partially mapped according to the analysis phase.

2 Main functionalities of MUMPS 4.7.1

We describe here the main functionalities of the solver MUMPS. The user should refer to Sections 4 and 5 for a complete description of the parameters that must be set or that are referred to in this Section. The variables mentioned in this section are components of a structure `mumps_par` of type `[SDCZ]MUMPS.STRUC` (see Section 3) and for the sake of clarity, we refer to them only by their component name. For example, we use `ICNTL` to refer to `mumps_par%ICNTL`.

2.1 Input matrix structure

MUMPS provides several possibilities for inputting the matrix. The selection is controlled by the parameters `ICNTL(5)` and `ICNTL(18)`.

The input matrix can be supplied in *elemental format* and must then be input centrally on the host (`ICNTL(5)=1` and `ICNTL(18)=0`). For full details see Section 4.6. Otherwise, it can be supplied in *assembled format* in coordinate form (`ICNTL(5)=0`), and, in this case, there are several possibilities (see Sections 4.5 and 4.7):

1. the matrix can be input centrally on the host processor (`ICNTL(18)=0`);
2. only the matrix structure is provided on the host for the analysis phase and the matrix entries are provided for the numerical factorization, distributed across the processors:
 - either according to a mapping supplied by the analysis (`ICNTL(18)=1`),
 - or according to a user determined mapping (`ICNTL(18)=2`);
3. it is also possible to distribute the matrix pattern and the entries in any distribution in local triplets (`ICNTL(18)=3`) for both analysis and factorization (recommended option for distributed entry).

By default the input matrix is considered in assembled format (`ICNTL(5)=0`) and input centrally on the host processor (`ICNTL(18)=0`).

2.2 Symmetric orderings

A range of orderings to preserve sparsity is available in the analysis phase. Most of them were introduced in Release 4.2 of the MUMPS package. The parameter `ICNTL(7)` is used to determine the ordering.

In addition to the approximate minimum degree ordering (AMD, [2]), an approximate minimum degree ordering with automatic quasi-dense row detection (QAMD, [1]), an approximate minimum fill-in ordering (AMF), an ordering where bottom-up strategies are used to build separators by Jürgen Schulze from University of Paderborn (PORD, [23]), and the METIS package from Univ. of Minnesota [22] are possible choices. When using the METIS package, only the METIS_NODEND hybrid ordering routine can be used.

A user-supplied ordering can also be provided and the pivot order must be set by the user in `PERM_IN` (see Section 4.9). Also, it should be noted that the logic that handles this case is different from the built-in orderings so that, for example, a different performance and different internal data structures are created by a run that generates an ordering and a separate one that feeds that same ordering array in as input.

If `ICNTL(7)=7`, the MUMPS package will automatically choose the ordering depending on the ordering packages installed, the type of the matrix (symmetric or unsymmetric), the size of the matrix and the number of processors available.

The default value of `ICNTL(7)` is 7.

2.3 Other pre-processing facilities

In addition to the symmetric orderings, MUMPS offers other pre-processing facilities: permuting to zero-free diagonal and prescaling.

Permutations to a zero-free diagonal can be applied to very unsymmetric matrices and can help reduce fill-in and arithmetic, see [14, 15]. This functionality is controlled by `ICNTL(6)`. For symmetric matrices this permutation can also be used to constrain the symmetric permutation (see `ICNTL(12)` option).

Prescaling of the input matrix can help reduce fill-in during factorization and can improve the numerical accuracy. A range of classical scalings are provided and can be automatically performed at the

beginning of the numerical factorization phase. This functionality is controlled by ICNTL(8). For some values of ICNTL(6) or ICNTL(12) the arrays COLSCA/ROWSCA can be allocated and built during the analysis phase (see Section 4.8). Symmetric indefinite matrices preprocessings, as described in [18], can be applied and are controlled by ICNTL(12).

2.4 Post-processing facilities

It has been shown [9] that with only two to three steps of iterative refinement the solution can often be significantly improved. Iterative refinement can be optionally performed after the solution step using the parameter ICNTL(10).

MUMPS also enables the user to perform classical error analysis based on the residuals (see the description of ICNTL(11) in Section 5). We calculate an estimate of the sparse backward error using the theory and metrics developed in [9]. We use the notation $\bar{\mathbf{x}}$ for the computed solution and a modulus sign on a vector or a matrix to indicate the vector or matrix obtained by replacing all entries by their moduli. The scaled residual

$$\frac{|\mathbf{b} - \mathbf{A}\bar{\mathbf{x}}|_i}{(|\mathbf{b}| + |\mathbf{A}| |\bar{\mathbf{x}}|)_i} \quad (1)$$

is computed for all equations except those for which the numerator is nonzero and the denominator is small. For all the exceptional equations,

$$\frac{|\mathbf{b} - \mathbf{A}\bar{\mathbf{x}}|_i}{(|\mathbf{A}| |\bar{\mathbf{x}}|)_i + \|\mathbf{A}_i\|_\infty \|\bar{\mathbf{x}}\|_\infty} \quad (2)$$

is used instead, where $\mathbf{A}_i$ is row i of $\mathbf{A}$. The largest scaled residual (1) is returned in RINFOG(7) and the largest scaled residual (2) is returned in RINFOG(8). If all equations are in category (1), zero is returned in RINFOG(8). The computed solution $\bar{\mathbf{x}}$ is the exact solution of the equation

$$(\mathbf{A} + \delta\mathbf{A})\mathbf{x} = (\mathbf{b} + \delta\mathbf{b}),$$

where

$$\delta\mathbf{A}_{ij} \leq \max(\text{RINFOG}(7), \text{RINFOG}(8)) |\mathbf{A}|_{ij},$$

and $\delta\mathbf{b}_i \leq \max(\text{RINFOG}(7) |\mathbf{b}|_i, \text{RINFOG}(8) \|\mathbf{A}_i\|_\infty \|\bar{\mathbf{x}}\|_\infty)$. Note that $\delta\mathbf{A}$ respects the sparsity of $\mathbf{A}$. An upper bound for the error in the solution is returned in RINFOG(9). Finally condition numbers cond_1 and cond_2 for the matrix are returned in RINFOG(10) and RINFOG(11), respectively, and

$$\frac{\|\delta\mathbf{x}\|}{\|\mathbf{x}\|} \leq \text{RINFOG}(7) \times \text{cond}_1 + \text{RINFOG}(8) \times \text{cond}_2.$$

2.5 Solving the transposed system

Given a sparse matrix $\mathbf{A}$, the system $\mathbf{A}\mathbf{X} = \mathbf{B}$ or $\mathbf{A}^T\mathbf{X} = \mathbf{B}$ can be solved during the solve stage, where $\mathbf{A}$ is square of order n and $\mathbf{X}$ and $\mathbf{B}$ are of order n by n_{rhs} . This is controlled by ICNTL(9).

2.6 Reduce/condense a problem on an interface (Schur complement and reduced/condensed right-hand side)

A Schur complement matrix (centralized or provided as 2D block cyclic matrix) can be returned to the user (see ICNTL(19), ICNTL(26) and Section 4.10). The user must specify the list of indices of the Schur matrix. MUMPS then provides both a partial factorization of the complete matrix and returns the assembled Schur matrix in user memory. The Schur matrix is considered as a full matrix. The partial factorization that builds the Schur matrix can also be used to solve linear systems associated with the “interior” variables (ICNTL(26)=0) and also to handle a reduced/condensed right-hand-side (ICNTL(26)=1,2) as described in the following discussion.

Let us consider a partitioned matrix (here with an unsymmetric matrix) where the variables of $\mathbf{A}_{2,2}$, specified by the user, correspond to the Schur variables and on which a partial factorization has been

performed. In the following, and only for the sake of clearness we have ordered last all variables belonging to the Schur.

$$\mathbf{A} = \begin{pmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} \end{pmatrix} = \begin{pmatrix} \mathbf{L}_{1,1} & 0 \\ \mathbf{L}_{2,1} & \mathbf{I} \end{pmatrix} \begin{pmatrix} \mathbf{U}_{1,1} & \mathbf{U}_{1,2} \\ 0 & \mathbf{S} \end{pmatrix} \quad (3)$$

Thus the Schur complement, as returned by MUMPS, is such that $\mathbf{S} = \mathbf{A}_{2,2} - \mathbf{A}_{2,1}\mathbf{A}_{1,1}^{-1}\mathbf{A}_{1,2}$.

ICNTL(26) can then be used during the solution phase to describe how this partial factorization can be used to solve $\mathbf{A}x = b$:

- **Compute a partial solution**:

If ICNTL(26)=0 then the solve is performed on the internal problem:

$$\mathbf{A}_{1,1}x_1 = b_1.$$

Entries in the right-hand side corresponding to indices from the Schur matrix need not be set on entry and they are explicitly set to zero on output.

- **Solve the complete system in three steps**:

$$\begin{pmatrix} \mathbf{L}_{1,1} & 0 \\ \mathbf{L}_{2,1} & \mathbf{I} \end{pmatrix} \begin{pmatrix} \mathbf{U}_{1,1} & \mathbf{U}_{1,2} \\ 0 & \mathbf{S} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} \quad (4)$$

1. **Reduction/condensation phase**:

One can compute with ICNTL(26)=1, the intermediate y vector, in which y_2 is often referred to as the reduced/condensed right-hand-side.

$$\begin{pmatrix} \mathbf{L}_{1,1} & 0 \\ \mathbf{L}_{2,1} & \mathbf{I} \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} \quad (5)$$

Then one has to solve

$$\begin{pmatrix} \mathbf{U}_{1,1} & \mathbf{U}_{1,2} \\ 0 & \mathbf{S} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} y_1 \\ y_2 \end{pmatrix} \quad (6)$$

2. **Using Schur matrix**:

The Schur matrix is an output of the factorisation phase. It is the responsibility of the user to compute x_2 such that $\mathbf{S}x_2 = y_2$.

3. **Expansion phase**:

Given x_2 and y_1 , option ICNTL(26)=2 of the solve phase can be used to compute x_1 . Note that the package uses y_1 computed (and stored in the `mumps` structure) during the first step (ICNTL(26)=1) and that the complete solution x is provided on output.

Note that the Schur complement could be considered as an element contribution to the interface block in a domain decomposition approach. MUMPS could then be used to solve this interface problem using the element entry functionality.

2.7 Arithmetic versions

Several versions of the package MUMPS are available: `REAL`, `DOUBLE PRECISION`, `COMPLEX`, and `DOUBLE COMPLEX`.

This document applies to all four precisions. In the following we use the conventions below:

1. the term **real** is used for `REAL` or `DOUBLE PRECISION`,
2. the term **complex** is used for `COMPLEX` or `DOUBLE COMPLEX`,
3. real version means either `REAL` or `DOUBLE PRECISION` version,
4. complex version means either `COMPLEX`, or `DOUBLE COMPLEX` version.

2.8 The working host processor

The analysis phase is performed on the host processor. This processor is the one with rank 0 in the communicator provided to MUMPS. By setting the variable PAR to 1 (see Section 4.3), MUMPS allows the host to participate in computations during the factorization and solve phases, just like any other processor. This allows MUMPS to run on a single processor and prevents the host processor being idle during the factorization and solve phases (as would be the case for PAR=0). We thus generally recommend using a working host processor (PAR=1).

The only case where it may be worth using PAR=0 is with a large centralized matrix on a purely distributed architecture with relatively small local memory: PAR=1 will lead to a memory imbalance because of the storage related to the initial matrix on the host.

2.9 Sequential version

It is possible to use MUMPS sequentially by limiting the number of processors to one, but the link phase still requires the MPI, BLACS, and ScaLAPACK libraries and the user program needs to make explicit calls to `MPI_INIT` and `MPI_FINALIZE`.

A purely sequential version of MUMPS is also available. For this, a special library is distributed that provides all external references needed by MUMPS for a sequential environment. MUMPS can thus be used in a simple sequential program, ignoring everything related to parallelism or MPI. Details on how to build a purely sequential version of MUMPS are available in the file `README` available in the MUMPS distribution. Note that for the sequential version, the component PAR must be set to 1 (see Section 4.3) and that the calling program should not make use of MPI.

2.10 Shared memory version

On networks of SMP nodes (multiprocessor nodes with a shared memory), a parallel shared memory BLAS library (also called multithread BLAS) is often provided by the manufacturer. Using shared memory BLAS (between 2 and 4 threads per MPI process) can be significantly more efficient than running with only MPI processes. For example on a computer with 2 SMP nodes and 16 processors per node, we advise to run using 16 MPI processes with 2 threads per MPI process.

2.11 Main changes between versions

MUMPS is an evolving package and new facilities and algorithms have been added that result in changes to the version number. Please refer to <http://www.enseeiht.fr/apo/MUMPS/faq.html> and <http://graal.ens-lyon.fr/MUMPS/faq.html> to an history of the main modifications where we also report the most frequently asked question about the solver. The `README` file with our distribution can be consulted for more details and for the dates of the releases.

3 Sequence in which routines are called

In the following, we use the notation `[SDCZ]MUMPS` to refer to `DMUMPS`, `SMUMPS`, `ZMUMPS` or `CMUMPS` for `REAL`, `DOUBLE PRECISION`, `COMPLEX` and `DOUBLE COMPLEX` versions, respectively. Similarly `[SDCZ]MUMPS_STRUC` refers to either `SMUMPS_STRUC`, `DMUMPS_STRUC`, `CMUMPS_STRUC`, or `ZMUMPS_STRUC`, and `[sdcz]mumps_struct.h` to `smumps_struct.h`, `dmumps_struct.h`, `cmumps_struct.h` or `zmumps_struct.h`.

In the Fortran 90 interface, there is a single user callable subroutine per precision, called `[SDCZ]MUMPS`, that has a single parameter `mumps_par` of Fortran 90 derived datatype `[SDCZ]MUMPS_STRUC` defined in `[sdcz]mumps_struct.h`. The interface is the same for the sequential version, only the compilation process and libraries need be changed. In the case of the parallel version, MPI must be initialized by the user before the first call to `[SDCZ]MUMPS` is made. The calling sequence for the `DOUBLE PRECISION` version may look as follows:

```
INCLUDE 'mpif.h'
INCLUDE 'dmumps_struct.h'
```

```

...
INTEGER IERR
TYPE (DMUMPS_STRUC) :: mumps_par
...
CALL MPI_INIT(IERR)      ! Not needed in purely sequential version
...
CALL DMUMPS( mumps_par )
...
CALL MPI_FINALIZE(IERR) ! Not needed in purely sequential version

```

For other precisions, `dmumps_struc.h` should be replaced by `smumps_struc.h`, `cmumps_struc.h`, or `zmumps_struc.h`, and the 'D' in `DMUMPS` and `DMUMPS_STRUC` by 'S', 'C' or 'Z'.

The variable `mumps_par` of datatype `[SDCZ]MUMPS_STRUC` holds all the data for the problem. It has many components, only some of which are of interest to the user. The other components are internal to the package. Some of the components must only be defined on the host. Others must be defined on all processors. The file `[sdcz]mumps_struc.h` defines the derived datatype and must always be included in the program that calls `MUMPS`. The file `[sdcz]mumps_root.h`, which is included in `[sdcz]mumps_struc.h`, must also be available at compilation time. Components of the structure `[SDCZ]MUMPS_STRUC` that are of interest to the user are shown in Figure 1.

The interface to `MUMPS` consists in calling the subroutine `[SDCZ]MUMPS` with the appropriate parameters set in `mumps_par`.

```

        INCLUDE '[sdcz]mumps_root.h'
        TYPE [SDCZ]MUMPS_STRUC
        SEQUENCE
C INPUT PARAMETERS
C *****
C Problem definition
C -----
C Solver (SYM=0 Unsymmetric, SYM=1 Sym. Positive Definite, SYM=2 General Symmetric)
C Type of parallelism (PAR=1 host working, PAR=0 host not working)
        INTEGER SYM, PAR, JOB
C Control parameters
C -----
        INTEGER ICNTL(40)

        real CNTL(15)

        INTEGER N ! Order of input matrix
C Assembled input matrix : User interface
C -----
        INTEGER NZ

        real/complex, DIMENSION(:), POINTER :: A

        INTEGER, DIMENSION(:), POINTER :: IRN, JCN
C Case of distributed matrix entry
C -----
        INTEGER NZ_loc
        INTEGER, DIMENSION(:), POINTER :: IRN_loc, JCN_loc

        real/complex, DIMENSION(:), POINTER :: A_LOC
C Unassembled input matrix: User interface
C -----
        INTEGER NELT
        INTEGER, DIMENSION(:), POINTER :: ELTPTR, ELTVAR

        real/complex, DIMENSION(:), POINTER :: A_ELT
C MPI Communicator
C -----
        INTEGER COMM
C Ordering and scaling, if given by user (optional)
C -----
        INTEGER, DIMENSION(:), POINTER :: PERM_IN

        real/complex, DIMENSION(:), POINTER :: COLSCA, ROWSCA
C INPUT/OUTPUT data
C *****
C RHS/SOL_loc : on input it holds the right-hand side
C on output it always holds the assembled solution
C -----
        real/complex, DIMENSION(:), POINTER :: RHS, REDRHS
        real/complex, DIMENSION(:), POINTER :: RHS_SPARSE

        INTEGER, DIMENSION(:), POINTER :: IRHS_SPARSE, IRHS_PTR
        INTEGER NRHS, LRHS, NZ_RHS, LSOL_LOC, LREDRHS

        real/complex, DIMENSION(:), POINTER :: SOL_LOC

        INTEGER, DIMENSION(:), POINTER :: ISOL_LOC
C OUTPUT data and Statistics
C *****
        INTEGER, DIMENSION(:), POINTER :: SYM_PERM, UNS_PERM
        INTEGER INFO(40)
        INTEGER INFOG(40) ! Global information (host only)

        real RINFO(20)
        real RINFOG(20) ! Global information (host only)
C Schur
        INTEGER SIZE_SCHUR, NPROW, NPCOL, MBLOCK, NBLOCK
        INTEGER SCHUR_MLOC, SCHUR_NLOC, SCHUR_LLD
        INTEGER, DIMENSION(:), POINTER :: LISTVAR_SCHUR

        real/complex, DIMENSION(:), POINTER :: SCHUR
C Mapping if provided by MUMPS
        INTEGER, DIMENSION(:), POINTER :: MAPPING
C Version number
        CHARACTER(LEN=80) VERSION_NUMBER

        END TYPE [SDCZ]MUMPS_STRUC

```

Figure 1: Main components of the structure [SDCZ]MUMPS_STRUC defined in [sdcz]mumps_struct.h. **real/complex** qualifies parameters that are real in the real version and complex in the complex version, whereas **real** is used for parameters that are always real, even in the complex version of MUMPS.

4 Input and output parameters

In this section, we describe the components of the variable `mumps_par` of datatype `[SDCZ]MUMPS_STRUC`. Those components define the arguments to MUMPS that must be set by the user, or that are returned to the user.

4.1 Version number

`mumps_par%VERSION_NUMBER` (string) is set by MUMPS to the version number of MUMPS after a call to the initialization phase (`JOB=-1`).

4.2 Control of the three main phases: Analysis, Factorization, Solve

`mumps_par%JOB` (integer) must be initialized by the user on all processors before a call to MUMPS. It controls the main action taken by MUMPS. It is not altered by MUMPS.

`JOB = -1` initializes an instance of the package. A call with `JOB = -1` must be performed before any other call to the package on the same instance. It sets default values for other components of `MUMPS_STRUC` (such as `ICNTL`, see below), which may then be altered before subsequent calls to MUMPS. Note that three components of the structure must always be set by the user (on all processors) before a call with `JOB = -1`. These are

- `mumps_par%COMM`,
- `mumps_par%SYM`, and
- `mumps_par%PAR`.

Note that, after a call to `JOB = -1`, the internal component `mumps_par%MYID` contains the rank of the calling processor in the communicator provided to MUMPS. Thus, the test “(`mumps_par%MYID == 0`)” may be used to identify the host processor (see Section 2.8).

Furthermore, the version number is returned in `mumps_par%VERSION_NUMBER` (see Section 4.1).

`JOB = -2` destroys an instance of the package. All data structures associated with the instance, except those provided by the user in `mumps_par`, are deallocated. It should be called by the user only when no further calls to MUMPS with this instance are required. It should be called before a further `JOB = -1` call with the same argument `mumps_par`.

`JOB=1` performs the analysis. In this phase, MUMPS chooses pivots from the diagonal using a selection criterion to preserve sparsity. It uses the pattern of $\mathbf{A} + \mathbf{A}^T$ but ignores numerical values. It subsequently constructs subsidiary information for the numerical factorization (a `JOB=2` call).

An option exists for the user to input the pivotal sequence (`ICNTL(7)=1`, see below) in which case only the necessary information for a `JOB=2` call will be generated.

The numerical values of the original matrix, `mumps_par%A`, must be provided by the user during the analysis phase only if `ICNTL(6)` is set to a value between 2 and 7. See `ICNTL(6)` in Section 5 for more details.

MUMPS uses the pattern of the matrix $\mathbf{A}$ input by the user. In the case of a *centralized matrix*, the following components of the structure defining the matrix pattern must be set by the user only on the host:

- `mumps_par%N`, `mumps_par%NZ`, `mumps_par%IRN`, and `mumps_par%JCN` if the user wishes to input the structure of the matrix in *assembled format* (`ICNTL(5)=0` and `ICNTL(18) ≠ 3`) (see Section 4.5),
- `mumps_par%N`, `mumps_par%NLT`, `mumps_par%ELTPTR`, and `mumps_par%ELTVAR` if the user wishes to input the matrix in *elemental format* (`ICNTL(5)=1`) (see Section 4.6).

These components should be passed unchanged when later calling the factorization (`JOB=2`) and solve (`JOB=3`) phases.

In the case of a *distributed assembled matrix* (see Section 4.7 for more details and options),

- If ICNTL(18) = 1 or 2, the previous requirements hold except that IRN and JCN are no longer required and need not be passed unchanged to the factorization phase.
- If ICNTL(18) = 3, the user should provide
 - mumps_par%N on the host
 - mumps_par%NZ_loc, mumps_par%IRN_loc and mumps_par%JCN_loc on all slave processors. Those should be passed unchanged to the factorization (JOB=2) and solve (JOB=3) phases.

A call to MUMPS with JOB=1 must be preceded by a call with JOB = -1 on the same instance.

JOB=2 performs the factorization. It uses the numerical values of the matrix **A** provided by the user and the information from the analysis phase (JOB=1) to factorize the matrix **A**.

If the matrix is centralized on the host (ICNTL(18)=0), the pattern of the matrix should be passed unchanged since the last call to the analysis phase (see JOB=1); the following components of the structure define the numerical values and must be set by the user (on the host only) before a call with JOB=2:

- mumps_par%A if the matrix is in assembled format (ICNTL(5)=0), or
- mumps_par%A_ELT if the matrix is in elemental format (ICNTL(5)=1).

If the initial matrix is distributed (ICNTL(5)=0 and ICNTL(18) $\neq$ 0), then the following components of the structure must be set by the user on all slave processors before a call with JOB=2:

- mumps_par%A_loc on all slave processors, and
- mumps_par%NZ_loc, mumps_par%IRN_loc and mumps_par%JCN_loc if ICNTL(18)=1 or 2. (For ICNTL(18)=3, NZ_loc, IRN_loc and JCN_loc have already been passed to the analysis step and must be passed unchanged.)

(See Sections 4.5, 4.6, and 4.7.)

The actual pivot sequence used during the factorization may differ slightly from the sequence returned by the analysis if the matrix **A** is not diagonally dominant.

An option exists for the user to input scaling vectors or let MUMPS compute such vectors automatically (in arrays COLSCA/ROWSCA, ICNTL(8) $\neq$ 0, see Section 4.8).

A call to MUMPS with JOB=2 must be preceded by a call with JOB=1 on the same instance.

JOB=3 performs the solution. It can also be used (see ICNTL(25)) to compute the null space basis provided that “null pivot row” detection (ICNTL(24)) was on and that the number of null pivots INFOG(28) was different from 0. It uses the right-hand side(s) **B** provided by the user and the factors generated by the factorization (JOB=2) to solve a system of equations $\mathbf{AX} = \mathbf{B}$ or $\mathbf{A}^T \mathbf{X} = \mathbf{B}$. The pattern and values of the matrix should be passed unchanged since the last call to the factorization phase (see JOB=2). The structure component mumps_par%RHS must be set by the user (on the host only) before a call with JOB=3. (See Section 4.12.)

A call to MUMPS with JOB=3 must be preceded by a call with JOB=2 (or JOB=4) on the same instance.

JOB=4 combines the actions of JOB=1 with those of JOB=2. It must be preceded by a call to MUMPS with JOB = -1 on the same instance.

JOB=5 combines the actions of JOB=2 and JOB=3. It must be preceded by a call to MUMPS with JOB=1 on the same instance.

JOB=6 combines the actions of calls with JOB=1, 2, and 3. It must be preceded by a call to MUMPS with JOB = -1 on the same instance.

Consecutive calls with JOB=2,3,5 on the same instance are possible.

4.3 Control of parallelism

mumps_par%COMM (integer) must be set by the user on all processors before the initialization phase (JOB = -1) and must not be changed. It must be set to a valid MPI communicator that will be used

for message passing inside MUMPS. It is not altered by MUMPS. The processor with rank 0 in this communicator is used by MUMPS as the **host** processor. Note that only the processors belonging to the communicator should call MUMPS.

mumps_par%**PAR** (integer) must be initialized by the user on all processors and is accessed by MUMPS only during the initialization phase (JOB = -1). It is not altered by MUMPS and its value is communicated internally to the other phases as required. Possible values for PAR are:

- 0 host is not involved in factorization/solve phases
- 1 host is involved in factorization/solve phases

Other values are treated as 1.

If PAR is set to 0, the host will only hold the initial problem, perform symbolic computations during the analysis phase, distribute data, and collect results from other processors. If set to 1, the host will also participate in the factorization and solve phases. If the initial problem is large and memory is an issue, PAR = 1 is not recommended if the matrix is centralized on processor 0 because this can lead to memory imbalance, with processor 0 having a larger memory load than the other processors. Note that setting PAR to 1, and using only 1 processor, leads to a sequential code.

4.4 Matrix type

mumps_par%**SYM** (integer) must be initialized by the user on all processors and is accessed by MUMPS only during the initialization phase (JOB = -1). It is not altered by MUMPS except for the complex version of MUMPS where SYM=1 is replaced by SYM=2 and structural symmetry is exploited up to the root. Its value is communicated internally to the other phases as required. Possible values for SYM are:

- 0 **A** is unsymmetric
- 1 **A** is symmetric positive definite
- 2 **A** is general symmetric

For the complex version, the value SYM=1 is currently treated as SYM=2. We do not have a version for Hermitian matrices in this release of MUMPS.

4.5 Centralized assembled matrix input: ICNTL(5)=0 and ICNTL(18)=0

mumps_par%**N** (integer), mumps_par%**NZ** (integer), mumps_par%**IRN** (integer array pointer, dimension NZ), mumps_par%**JCN** (integer array pointer, dimension NZ), and mumps_par%**A** (**real/complex** array pointer, dimension NZ) hold the matrix in assembled format. These components should be set by the user only on the host and only when ICNTL(5)=0 and ICNTL(18)=0; they are not modified by the package.

- **N** is the order of the matrix **A**, $N > 0$. It is not altered by MUMPS.
- **NZ** is the number of entries being input, $NZ > 0$. It is not altered by MUMPS.
- **IRN**, **JCN** are integer arrays of length **NZ** containing the row and column indices, respectively, for the matrix entries.
- **A** is a **real** (**complex** in the complex version) array of length **NZ**. The user must set **A(k)** to the value of the entry in row **IRN(k)** and column **JCN(k)** of the matrix. **A** is accessed when **JOB=1** only when **ICNTL(6) $\neq$ 0**. Duplicate entries are summed and any with **IRN(k)** or **JCN(k)** out-of-range are ignored.
 Note that, in the case of the symmetric solver, a diagonal nonzero a_{ii} is held as **A(k)= a_{ii}** , **IRN(k)=JCN(k)= i** , and a pair of off-diagonal nonzeros $a_{ij} = a_{ji}$ is held as **A(k)= a_{ij}** and **IRN(k)= i** , **JCN(k)= j** or vice-versa. Again, duplicate entries are summed and entries with **IRN(k)** or **JCN(k)** out-of-range are ignored.

The components **N**, **NZ**, **IRN**, and **JCN** describe the pattern of the matrix and must be set by the user before the analysis phase (**JOB=1**). Component **A** must be set before the factorization phase (**JOB=2**).

4.6 Element matrix input: ICNTL(5)=1 and ICNTL(18)=0

mumps_par%N (integer), mumps_par%NELT (integer), mumps_par%ELTPTR (integer array pointer, dimension NELT+1), mumps_par%ELTVAR (integer array pointer, dimension ELTPTR(NELT+1) – 1), and mumps_par%A_ELTVAR (**real/complex** array pointer) hold the matrix in elemental format. These components should be set by the user only on the host and only when ICNTL(5)=1:

- N is the order of the matrix **A**, $N > 0$. It is not altered by MUMPS.
- NELT is the number of elements being input, $NELT > 0$. It is not altered by MUMPS.
- ELTPTR is an integer array of length NELT+1. ELTPTR(j) points to the position in ELTVAR of the first variable in element j, and ELTPTR(NELT+1) must be set to the position after the last variable of the last element. Note that ELTPTR(1) should be equal to 1. ELTPTR is not altered by MUMPS.
- ELTVAR is an integer array of length ELTPTR(NELT+1) – 1 and must be set to the lists of variables of the elements. It is not altered by MUMPS. Those for element j are stored in positions ELTPTR(j), ..., ELTPTR(j+1)–1. Out-of-range variables are ignored.
- A_ELTVAR is a **real (complex in the complex version)** array. If N_p denotes ELTPTR(p+1)–ELTPTR(p), then the values for element j are stored in positions $K_j + 1, \dots, K_j + L_j$, where
 - $K_j = \sum_{p=1}^{j-1} N_p^2$, and $L_j = N_j^2$ in the unsymmetric case ($SYM = 0$)
 - $K_j = \sum_{p=1}^{j-1} (N_p \cdot (N_p + 1))/2$, and $L_j = (N_j \cdot (N_j + 1))/2$ in the symmetric case ($SYM \neq 0$). Only the lower triangular part is stored.

Values within each element are stored column-wise. Values corresponding to out-of-range variables are ignored and values corresponding to duplicate variables within an element are summed. A_ELTVAR is not accessed when JOB = 1. Note that, although the elemental matrix may be symmetric or unsymmetric in value, its structure is always symmetric.

The components N, NELT, ELTPTR, and ELTVAR describe the pattern of the matrix and must be set by the user before the analysis phase (JOB=1). Component A_ELTVAR must be set before the factorization phase (JOB=2). Note that, in the current release of the package, the element entry must be centralized on the host.

4.7 Distributed assembled matrix input: ICNTL(5)=0 and ICNTL(18)≠0

When the matrix is in assembled form (ICNTL(5)=0), we offer several options, defined by the control parameter ICNTL(18) described in Section 5. The following components of the structure define the distributed assembled matrix input. They are valid for nonzero values of ICNTL(18), otherwise the user should refer to Section 4.5.

mumps_par%N (integer), mumps_par%NZ (integer), mumps_par%IRN (integer array pointer, dimension NZ), mumps_par%JCN (integer array pointer, dimension NZ), mumps_par%IRN_loc (integer array pointer, dimension NZ_loc), mumps_par%JCN_loc (integer array pointer, dimension NZ_loc), mumps_par%A_loc (**real/complex** array pointer, dimension NZ_loc), and mumps_par%MAPPING (integer array, dimension NZ).

- N is the order of the matrix **A**, $N > 0$. It must be set on the host before analysis. It is not altered by MUMPS.
- NZ is the number of entries being input in the definition of **A**, $NZ > 0$. It must be defined on the host before analysis if ICNTL(18) = 1, or 2.
- IRN, JCN are integer arrays of length NZ containing the row and column indices, respectively, for the matrix entries. They must be defined on the host before analysis if ICNTL(18) = 1, or 2. They can be deallocated by the user just after the analysis.
- NZ_loc is the number of entries local to a processor. It must be defined on all processors in the case of the working host model of parallelism (PAR=1), and on all processors except the host in the case of the non-working host model of parallelism (PAR=0), before analysis if ICNTL(18) = 3, and before factorization if ICNTL(18) = 1 or 2.

- `IRN_loc`, `JCN_loc` are integer arrays of length `NZ_loc` containing the row and column indices, respectively, for the matrix entries. They must be defined on all processors if `PAR=1`, and on all processors except the host if `PAR=0`, before analysis if `ICNTL(18) = 3`, and before factorization if `ICNTL(18) = 1` or `2`.
- `A_loc` is a **real (complex in the complex version)** array of dimension `NZ_loc` that must be defined before the factorization phase (`JOB=2`) on all processors if `PAR = 1`, and on all processors except the host if `PAR = 0`. The user must set `A_loc(k)` to the value in row `IRN_loc(k)` and column `JCN_loc(k)`.
- `MAPPING` is an integer array of size `NZ` which is returned by MUMPS on the host after the analysis phase as an indication of a preferred mapping if `ICNTL(18) = 1`. In that case, `MAPPING(i) = IPROC` means that entry `IRN(i)`, `JCN(i)` should be provided on processor with rank `IPROC` in the MUMPS communicator.

We recommend the use of options `ICNTL(18)=2` or `3` because they are the simplest and most flexible options. Furthermore, those options (`2` or `3`) are in general almost as efficient as the more sophisticated (but more complicated for the user) option `ICNTL(18)=1`.

4.8 Scaling

`mumps_par%COLSCA`, `mumps_par%ROWSCA` (double precision array pointers, dimension `N`) are optional scaling arrays required only by the host. If a scaling is provided by the user (`ICNTL(8) = -1`), these arrays must be allocated and initialized by the user on the host, before a call to the factorization phase (`JOB=2`). They might also be automatically allocated and computed by the package during analysis (if `ICNTL(6)=5` or `6`), in which case `ICNTL(8) = -2` will be set by the package during analysis and should be passed unchanged to the solve phase (`JOB=3`).

4.9 Given ordering: ICNTL(7)=1

`mumps_par%PERM_IN` (integer array pointer, dimension `N`) must be allocated and initialized by the user on the host if `ICNTL(7)=1`. It is accessed during the analysis (`JOB=1`) and `PERM_IN(i)`, $i=1, \dots, N$ must hold the position of variable i in the pivot order. Note that, even when the ordering is provided by the user, the analysis must still be performed before numerical factorization.

4.10 Schur complement with reduced (or condensed) right-hand side: ICNTL(19) and ICNTL(26)

`mumps_par%SIZE_SCHUR` (integer) must be initialized on the host to the number of variables defining the Schur complement if `ICNTL(19) = 1, 2, or 3`. It is accessed during the analysis phase and should be passed unchanged to the factorization and solve phases.

`mumps_par%LISTVAR_SCHUR` (integer array pointer, dimension `mumps_par%SIZE_SCHUR`) must be allocated and initialized by the user on the host if `ICNTL(19) = 1, 2 or 3`. It is not altered by MUMPS. It is accessed during analysis (`JOB=1`) and `LISTVAR_SCHUR(i)`, $i=1, \dots, \text{SIZE_SCHUR}$ must hold the i^{th} variable of the Schur complement matrix.

Centralized Schur complement (ICNTL(19)=1)

`mumps_par%SCHUR` is a **real (complex in the complex version)** 1-dimensional pointer array that should point to size `SIZE_SCHUR × SIZE_SCHUR` locations in memory. It must be allocated by the user on the host (independently of the value of `mumps_par%PAR`) before the factorization phase. On exit, it holds the Schur complement matrix. On output from the factorization phase, and on the host node, the 1-dimensional pointer array `SCHUR` of length `SIZE_SCHUR*SIZE_SCHUR` holds the (dense) Schur matrix of order `SIZE_SCHUR`. Note that the order of the indices in the Schur matrix is identical to the order provided by the user in `LISTVAR_SCHUR` and that the Schur matrix is stored **by rows**. If the matrix is symmetric then only the lower triangular part of the Schur matrix is provided (**by rows**) and the upper part is not significant. (This can also be viewed as the upper triangular part stored by columns in which case the lower part is not defined.)

Distributed Schur complement (ICNTL(19)=2 or 3)

For symmetric matrices, the value of ICNTL(19) controls whether only the lower part (ICNTL(19) = 2) or the complete matrix (ICNTL(19) = 3) is generated. We always provide the complete matrix for unsymmetric matrices so either value for ICNTL(19) has the same effect.

If ICNTL(19)=2 or 3, the following parameters should be defined on the host on entry to the analysis phase:

mumps_par%NPROW, mumps_par%NPCOL, mumps_par%MBLOCK, and mumps_par%NBLOCK are integers corresponding to the characteristics of a 2D block cyclic grid of processors. They should be defined on the host before a call to the analysis phase. If any of these quantities is smaller than or equal to zero or has not been defined by the user, or if $\text{NPROW} \times \text{NPCOL}$ is larger than the number of slave processors available (total number of processors if mumps_par%PAR=1, total number of processors minus 1 if mumps_par%PAR=0), then a grid shape will be computed by the analysis phase of MUMPS and NPROW, NPCOL, MBLOCK, NBLOCK will be overwritten on exit from the analysis phase. Please refer to [10] (for example) for more details on the notion of grid of processors and on 2D block cyclic distributions. We briefly describe the meaning of the four above parameters here:

- NPROW is the number of processors in a row of the process grid,
- NPCOL is the number of processors in a column of the process grid,
- MBLOCK is the blocking factor used to distribute the rows of the Schur complement,
- NBLOCK is the blocking factor used to distribute the columns of the Schur complement.

As in ScaLAPACK, we use a row-major process grid of processors, that is, process ranks (as provided to MUMPS in the MPI communicator) are consecutive in a row of the process grid. NPROW, NPCOL, MBLOCK and NBLOCK should be passed unchanged from the analysis phase to the factorization phase.

On exit from the analysis phase, the following two components are set by MUMPS on the first $\text{NPROW} \times \text{NPCOL}$ slave processors (the host is excluded if PAR=0 and the processors with largest MPI ranks in the communicator provided to MUMPS may not be part of the grid of processors).

mumps_par%SCHUR_MLOC is an integer giving the number of rows of the local Schur complement matrix on the concerned processor. It is equal to $\text{MAX}(1, \text{NUMROC}(\text{SIZE_SCHUR}, \text{MBLOCK}, \text{myrow}, 0, \text{NPROW}))$, where

- NUMROC is an INTEGER function defined in most ScaLAPACK implementations (also used internally by the MUMPS package),
- SIZE_SCHUR, MBLOCK, NPROW have been defined earlier, and
- myrow is defined as follows:
Let myid be the rank of the calling process in the communicator COMM provided to MUMPS. (myid can be returned by the MPI routine MPI_COMM_RANK.)
 - if PAR = 1 myrow is equal to $\text{myid} / \text{NPCOL}$,
 - if PAR = 0 myrow is equal to $(\text{myid} - 1) / \text{NPCOL}$.

Note that an upperbound of the minimum value of leading dimension (SCHUR_LLD defined below) is equal to $((\text{SIZE_SCHUR} + \text{MBLOCK} - 1) / \text{MBLOCK} + \text{NPROW} - 1) / \text{NPROW} * \text{MBLOCK}$.

mumps_par%SCHUR_NLOC is an integer giving the number of columns of the local Schur complement matrix on the concerned processor. It is equal to $\text{NUMROC}(\text{SIZE_SCHUR}, \text{NBLOCK}, \text{mycol}, 0, \text{NPCOL})$, where

- SIZE_SCHUR, NBLOCK, NPCOL have been defined earlier, and
- mycol is defined as follows:
Let myid be the rank of the calling process in the communicator COMM provided to MUMPS. (myid can be returned by the MPI routine MPI_COMM_RANK.)
 - if PAR = 1 mycol is equal to $\text{MOD}(\text{myid}, \text{NPCOL})$,
 - if PAR = 0 mycol is equal to $\text{MOD}(\text{myid} - 1, \text{NPCOL})$.

On entry to the factorization phase (JOB = 2), SCHUR_LLD should be defined by the user and SCHUR should be allocated by the user on the NPROW $\times$ NPCOL first slave processors (the host is excluded if PAR=0 and the processors with largest MPI ranks in the communicator provided to MUMPS may not be part of the grid of processors).

mumps_par%**SCHUR_LLD** is an integer defining the leading dimension of the local Schur complement matrix. It should be larger or equal to the local number of rows of that matrix, SCHUR_MLOC (as returned by MUMPS on exit from the analysis phase on the processors that participate in the computation of the Schur). SCHUR_LLD is not modified by MUMPS.

mumps_par%**SCHUR** is a **real (complex in the complex version)** one-dimensional pointer array that should be allocated by the user before a call to the factorization phase. Its size should be at least equal to SCHUR_LLD $\times$ (SCHUR_NLOC - 1) + SCHUR_MLOC, where SCHUR_MLOC, SCHUR_NLOC, and SCHUR_LLD have been defined above. On exit to the factorization phase, the pointer array SCHUR contains the Schur complement, stored by columns, in the format corresponding to the 2D cyclic grid of NPROW $\times$ NPCOL processors, with block sizes MBLOCK and NBLOCK, and local leading dimensions SCHUR_LLD.

The Schur complement is stored by columns. Note that setting NPCOL $\times$ NPROW = 1 will centralize the Schur complement matrix, *stored by columns* (instead of by rows as in the ICNTL(19)=1 option). It will then be available on the host node if PAR=1, and on the node with MPI identifier 1 (first working slave processor) if PAR=0.

If ICNTL(19)=2 and the Schur is symmetric (SYM=1 or 2), only the lower triangle is provided, stored by columns.

If ICNTL(19)=3 and the Schur is symmetric (SYM=1 or 2), then both the lower and upper triangles are provided, stored by columns. Note that if ICNTL(19)=3, then the constraint mumps_par%MBLOCK = mumps_par%NBLOCK should hold.

(For unsymmetric matrices, ICNTL(19)=2 and ICNTL(19)=3 have the same effect.)

Using partial factorization during solution phase (ICNTL(26)= 0, 1 or 2)

As explained in Section 2.6, when a Schur complement has been computed during the factorization phase, then either the solution phase computes a solution on the internal problem (ICNTL(26)=0, see control parameter ICNTL(26)), or the complete problem can use a reduced right-hand side to build the solution of the problem on the Schur variables (ICNTL(26)=1 and ICNTL(26)=2).

If ICNTL(26)=1 or 2, then the following parameters must be defined on the host on entry to the solution step:

mumps_par%**LREDRHS** is an integer defining the leading dimension of the reduced right-hand side, REDRHS. It must be larger or equal to SIZE_SCHUR, the size of the Schur complement.

mumps_par%**REDRHS** is a **real (complex in the complex version)** one-dimensional pointer array that should be allocated by the user before entering the solution phase. Its size should be at least equal to LREDRHS $\times$ (NRHS-1) + SIZE_SCHUR. If ICNTL(26)=1, then on exit from the solution phase, REDRHS(i+(k-1)*LREDRHS), i=1, ..., SIZE_SCHUR, k=1, ..., NRHS will hold the reduced right-hand side. If ICNTL(26)=2, then REDRHS(i+(k-1)*LREDRHS), i=1, ..., SIZE_SCHUR, k=1, ..., NRHS must be set (on entry to the solution phase) to the solution on the Schur variables. In that case (ie, ICNTL(26)=2), it is not altered by MUMPS.

4.11 Workspace parameters

mumps_par%**MAXIS** and mumps_par%**MAXS** (integers) are defined, for each processor, as the size of the integer and the real (complex for the complex version) workspaces respectively required for factorization and/or solve. On return from analysis (JOB = 1), INFO(7) and INFO(8) return the minimum values for MAXIS and MAXS, respectively, to the user. If the user has reason to believe that significant numerical pivoting will be required, it may be desirable to choose a higher value for MAXIS (or MAXS) than output from the analysis, or to increase the value of ICNTL(14). At the beginning of the factorization, MAXIS and MAXS are set to the maximum of estimates based on analysis phase data (but including the

memory relaxation resulting from the value of ICNTL(14) provided to the factorization) and the values supplied by the user. An integer array IS of size MAXIS and a real (complex in the complex version) array S of size MAXS are then dynamically allocated and used during the factorization and solve phases to hold the factors and contribution blocks.

4.12 Right-hand side and solution vectors/matrices

The formats of the right-hand side and of the solution are controlled by ICNTL(20) and ICNTL(21), respectively.

Centralized dense right-hand side (ICNTL(20)=0) and/or centralized dense solution (ICNTL(21)=0)

If ICNTL(20)=0 or ICNTL(21)=0, the following should be defined on the host.

mumps_par%**RHS** (**real/complex** array pointer, dimension NRHS×LRHS) is a **real (complex** in the complex version) array that should be allocated by the user on the host before a call to MUMPS with JOB= 3, 5, or 6.

On entry, if ICNTL(20)=0, RHS(i+(k-1)×LRHS) must hold the i-th component of kth right-hand side vector of the equations being solved.

On exit, if ICNTL(21)=0, then RHS(i+(k-1)×LRHS) will hold the i-th component of the kth solution vector.

mumps_par%**NRHS** (integer) is an optional parameter that is significant on the host before a call to MUMPS with JOB = 3, 5, or 6. If set, it should hold the number of right-hand side vectors. If not set, the value 1 is assumed to ensure backward compatibility of the MUMPS interface with versions prior to 4.3.3. Note that if NRHS > 1, then functionalities related to iterative refinement and error analysis (see ICNTL(10) and ICNTL(11)) are currently disabled.

mumps_par%**LRHS** (integer) is an optional parameter that is significant on the host before a call to MUMPS with JOB=3, 5, or 6. If NRHS is provided, LRHS should then hold the leading dimension of the array RHS. Note that in that case, LRHS should be greater than or equal to N.

Sparse right-hand side (ICNTL(20)=1)

If ICNTL(20)=1, the following input parameters should be defined on the host only before a call to MUMPS with JOB=3, 5, or 6:

mumps_par%**NZ_RHS** (integer) should hold the number of non-zeros in all the right-hand side vectors.

mumps_par%**NRHS** (integer), if set, should hold the number of right-hand side vectors. If not set, the value 1 is assumed.

mumps_par%**RHS_SPARSE** (**real/complex** array pointer, dimension NZ_RHS) should hold the numerical values of the non-zero inputs of each right-hand side vector. See also IRHS_PTR below.

mumps_par%**IRHS_SPARSE** (integer array pointer, dimension NZ_RHS) should hold the indices of the variables of the non-zero inputs of each right-hand side vector.

mumps_par%**IRHS_PTR** is an integer array pointer of dimension NRHS+1. IRHS_PTR is such that the i-th right-hand side vector is defined by its non-zero row indices IRHS_SPARSE(IRHS_PTR(i)...IRHS_PTR(i+1)-1) and the corresponding numerical values RHS_SPARSE(IRHS_PTR(i)...IRHS_PTR(i+1)-1). Note that IRHS_PTR(1)=1 and IRHS_PTR(NRHS+1)=NZ_RHS+1.

Note that, if the right-hand side is sparse and the solution is centralized (ICNTL(21)=0), then mumps_par%**RHS** should still be allocated on the host, as explained in the previous section. On exit from a call to MUMPS with JOB=3, 5, or 6, it will hold the centralized solution.

Distributed solution (ICNTL(21)=1)

On some networks with low bandwidth, and especially when there are many right-hand side vectors, centralizing the solution on the host processor might be a costly operation in the solution phase from MUMPS. If this is critical to the user, this functionality allows the solution to be left distributed over the processors. The solution should then be exploited in its distributed form by the user application.

mumps_par%**SOL_LOC** is a **real/complex** array pointer, of dimension $\text{LSOL_LOC} \times \text{NRHS}$ (where NRHS corresponds to the value provided in mumps_par%NRHS on the host), that should be allocated by the user before the solve phase (JOB=3) on all processors in the case of the working host model of parallelism (PAR=1), and on all processors except the host in the case of the non-working host model of parallelism (PAR=0). Its leading dimension LSOL_LOC should be larger than or equal to INFO(23), where INFO(23) has the value returned by MUMPS on exit from the factorization phase. On exit from the solve phase, $\text{SOL_LOC}(i+(k-1) \times \text{LSOL_LOC})$ will contain the value corresponding to variable ISOL_LOC(i) in the k^{th} solution vector.

mumps_par%**LSOL_LOC** (integer). LSOL_LOC must be set to the leading dimension of SOL_LOC (see above) and should be larger than or equal to INFO(23), where INFO(23) has the value returned by MUMPS on exit from the factorization phase.

mumps_par%**ISOL_LOC** (integer array pointer, dimension INFO(23)) ISOL_LOC should be allocated by the user before the solve phase (JOB=3) on all processors in the case of the working host model of parallelism (PAR=1), and on all processors except the host in the case of the non-working host model of parallelism (PAR=0). ISOL_LOC should be of size at least INFO(23), where INFO(23) has the value returned by MUMPS on exit from the factorization phase. On exit from the solve phase, ISOL_LOC(i) contains the index of the variables for which the solution (in SOL_LOC) is available on the local processor. Note that if successive calls to the solve phase (JOB=3) are performed for a given matrix, ISOL_LOC will have the same contents for each of these calls.

Note that if the solution is kept distributed, then functionalities related to error analysis and iterative refinement (see ICNTL(10) and ICNTL(11)) are currently not available.

5 Control parameters

On exit from the initialization call (JOB = -1), the control parameters are set to default values. If the user wishes to use values other than the defaults, the corresponding entries in mumps_par%ICNTL and mumps_par%CNTL should be reset after this initial call and before the call in which they are used.

mumps_par%**ICNTL** is an integer array of dimension 40.

ICNTL(1) is the output stream for error messages. If it is negative or zero, these messages will be suppressed. Default value is 6.

ICNTL(2) is the output stream for diagnostic printing, statistics, and warning messages. If it is negative or zero, these messages will be suppressed. Default value is 0.

ICNTL(3) is the output stream for global information, collected on the host. If it is negative or zero, these messages will be suppressed. Default value is 6.

ICNTL(4) is the level of printing for error, warning, and diagnostic messages. Maximum value is 4 and default value is 2 (errors and warnings printed). Possible values are

- ≤ 0 : No messages output.
- 1 : Only error messages printed.
- 2 : Errors, warnings, and main statistics printed.
- 3 : Errors and warnings and terse diagnostics (only first ten entries of arrays) printed.
- 4 : Errors and warnings and information on input and output parameters printed.

ICNTL(5) has default value 0 and is only accessed by the host and only during the analysis phase. If ICNTL(5) = 0, the input matrix must be given in assembled format in the structure components N,

NZ, IRN, JCN, and A (or NZ_loc, IRN_loc, JCN_loc, A_loc, see Section 4.7). If ICNTL(5) = 1, the input matrix must be given

N, NELT, ELTPTR, ELTVAR, and A.ELT.

ICNTL(6) has default value 7 (automatic choice done by the package) and is used to control an option for permuting and scaling the matrix. It is only accessed by the host and only during the analysis phase. For unsymmetric matrices, if ICNTL(6)=1, 2, 3, 4, 5, 6 a column permutation (based on weighted bipartite matching algorithms described in [14, 15]) is applied to the original matrix to get a zero-free diagonal. For symmetric matrices, if ICNTL(6)=1, 2, 3, 4, 5, 6 a set of recommended 1×1 and 2×2 pivots is computed (see [18] for more details) from the computed column permutation. Possible values of ICNTL(6) are:

- 0 : No column permutation is computed.
- 1 : The permuted matrix has as many entries on its diagonal possible. The values on the diagonal are of arbitrary size.
- 2 : The smallest value on the diagonal of the permuted matrix is maximized.
- 3 : Variant of option 2 with different performance.
- 4 : The sum of the diagonal entries of the permuted matrix is maximized.
- 5 : The product of the diagonal entries of the permuted matrix is maximized. Vectors are also computed (and stored in COLSCA and ROWSCA, only if ICNTL(8) is set to 7) to scale the permuted matrix so that the nonzero diagonal entries in the permuted matrix are one in absolute value and all the off-diagonal entries are less than or equal to one in absolute value.
- 6 : Similar to 5 but with a different algorithm.
- 7 : Based on the structural symmetry of the input matrix and on the availability of the numerical values, the value of ICNTL(6) is automatically chosen by the software.

Other values are treated as 0.

Except for ICNTL(6)=0 or 1, the numerical values of the original matrix, mumps_par%A, must be provided by the user during the analysis phase. If the matrix is symmetric positive definite (SYM = 1), or in elemental format (ICNTL(5)=1), or the ordering is provided by the user (ICNTL(7)=1), or the Schur option (ICNTL(19) = 1, 2, or 3) is required, or the matrix is initially distributed (ICNTL(18) ≠ 0), then ICNTL(6) is treated as 0.

On unsymmetric matrices (SYM = 0), the user is advised to set ICNTL(6) to a nonzero value when the matrix is very unsymmetric in structure. On output from the analysis phase, when the column permutation is not the identity, the pointer mumps_par%UNS_PERM (internal data valid until a call to MUMPS with JOB=-2) provides access to the permutation. (The column permutation is such that entry $a_{i,perm(i)}$ is on the diagonal of the permuted matrix.) Otherwise, the pointer is unassociated.

On general symmetric matrices (SYM = 2), we advise either to let MUMPS select the strategy (ICNTL(6) = 7) or to set ICNTL(6) = 5 if the user knows that the matrix is for example an augmented system (which is a system with a large zero diagonal block). On output from the analysis the pointer mumps_par%UNS_PERM is unassociated.

On output from the analysis phase, INFOG(23) holds the value of ICNTL(6) that was effectively used.

ICNTL(7) has default value 7 and is only accessed by the host and only during the analysis phase. It determines the pivot order to be used for the factorization. Note that, even when the ordering is provided by the user, the analysis must be performed before numerical factorization. In exceptional cases, ICNTL(7) may be modified by MUMPS when the ordering is not compatible with the value of ICNTL(12). Possible values are:

- 0 : Approximate Minimum Degree (AMD) [2] is used,
- 1 : the pivot order should be set by the user in PERM_IN. In this case, PERM_IN(i), (i=1, ... N) holds the position of variable i in the pivot order.
- 2 : Approximate Minimum Fill (AMF) is used,
- 3 : Not available in the current version (treated as 7).

- 4 : PORD¹ [23] is used,
- 5 : the METIS² [22] routine METIS_NODEND is used,
- 6 : Approximate Minimum Degree with automatic quasi-dense row detection (QAMD) is used.
- 7 : Automatic choice by the software during analysis phase. This choice will depend on the ordering packages made available, on the matrix (type and size), and on the number of processors.

Other values are treated as 7. Currently, options 4 and 5 are only available if the corresponding packages are installed (see comments in the Makefiles to let MUMPS know about them). If the packages are not installed then options 4 and 5 are treated as 7.

- If the user asks for a Schur complement matrix and the matrix is assembled then only options 0, 1, 5 and 7 are currently available. Other options are treated as 7 which will (currently) be treated as 5 if the METIS package has been installed (strongly recommended) or 0 (AMD) otherwise.
- For elemental matrices (ICNTL(5)=1), only options 0, 1, 5 and 7 are available, with option 7 leading to an automatic choice between AMD and METIS (options 0 or 5); other values are treated as 7. Furthermore, if the user asks for a Schur complement matrix, only options 0, 1 and 7 are currently available. Other options are treated as 7 which will (currently) be treated as 0 (AMD).

Generally, with the automatic choice corresponding to ICNTL(7)=7, the option chosen by the package depends on the ordering packages installed, the type of matrix (symmetric or unsymmetric), the size of the matrix and the number of processors.

For matrices with relatively dense rows, we highly recommend option 6 which may significantly reduce the time for analysis.

On output, the pointer `mumps_par%SYM_PERM` provides access to the symmetric permutation that is effectively used by the MUMPS package, and `INFOG(7)` to the ordering option that was effectively used. (`mumps_par%SYM_PERM(i)`, ($i=1, \dots, N$) holds the position of variable i in the pivot order.)

ICNTL(8) has default value 7. It is used to describe the scaling strategy and is only accessed by the host.

On entry to the analysis phase, if ICNTL(8) = 7, then an automatic choice of the scaling option is performed during the analysis and ICNTL(8) is modified accordingly. In particular, if ICNTL(8) is set to -2 by the user or reset to -2 by the package during the analysis, scaling arrays are computed internally and will be ready to be used by the factorization phase.

On entry to the factorization phase, if ICNTL(8) = -1, scaling vectors must be provided in COLSCA and ROWSCA by the user, who is then responsible for allocating and freeing them, if ICNTL(8) = -2, scaling vectors must be provided in COLSCA and ROWSCA by the package (see previous paragraph). If ICNTL(8) = 0, no scaling is performed, and arrays COLSCA/ROWSCA are not used. If ICNTL(8) > 0, the scaling arrays COLSCA/ROWSCA are allocated and computed by the package during the factorization phase.

Possible values of ICNTL(8) are listed below:

- -2: Scaling computed during analysis (see [14, 15] for the unsymmetric case and [18] for the symmetric case).
- -1: Scaling provided on entry to the numerical factorization phase,
- 0 : No scaling applied/computed.
- 1 : Diagonal scaling,
- 2 : Scaling based on [11],
- 3 : Column scaling,
- 4 : Row and column scaling,

¹Distributed within MUMPS by permission of J. Schulze (University of Paderborn).

²See <http://glaros.dtc.umn.edu/gkhome/metis/metis/overview> to obtain a copy.

- 5 : Scaling based on [11] followed by column scaling,
- 6 : Scaling based on [11] followed by row and column scaling.
- 7 (analysis only) : Automatic choice of scaling value done during analysis.

If the input matrix is symmetric ($\text{SYM} \neq 0$), then only options -2 , -1 , 0 , 1 and 7 are allowed and other options are treated as 0 ; if $\text{ICNTL}(8) = -1$, the user should ensure that the array ROWSCA is equal to the array COLSCA. If the input matrix is in elemental format ($\text{ICNTL}(5) = 1$), then only options -1 and 0 are allowed and other options are treated as 0 . If the initial matrix is distributed ($\text{ICNTL}(18) \neq 0$ and $\text{ICNTL}(5) = 0$) then the value of $\text{ICNTL}(8)$ is ignored and no scaling is applied. If $\text{ICNTL}(8) = -2$ then the user has to provide the numerical value (in `mumps_par%A`) on entry to the analysis.

$\text{ICNTL}(9)$ has default value 1 and is only accessed by the host during the solve phase. If $\text{ICNTL}(9) = 1$, $\mathbf{Ax} = \mathbf{b}$ is solved, otherwise, $\mathbf{A}^T \mathbf{x} = \mathbf{b}$ is solved.

$\text{ICNTL}(10)$ has default value 0 and is only accessed by the host during the solve phase. If $\text{NRHS} = 1$, then $\text{ICNTL}(10)$ corresponds to the maximum number of steps of iterative refinement. If $\text{ICNTL}(10) \leq 0$, iterative refinement is not performed.

In the current version, if $\text{ICNTL}(21)=1$ (solution kept distributed) or $\text{NRHS} > 1$, then iterative refinement is not performed and $\text{ICNTL}(10)$ is treated as 0 .

$\text{ICNTL}(11)$ has default value 0 and is only accessed by the host and only during the solve phase. A positive value will return statistics related to the linear system solved ($\mathbf{Ax} = \mathbf{b}$ or $\mathbf{A}^T \mathbf{x} = \mathbf{b}$ depending on the value of $\text{ICNTL}(9)$): the infinite norm of the input matrix, the computed solution, and the scaled residual in $\text{RINFOG}(4)$ to $\text{RINFOG}(6)$, respectively, a backward error estimate in $\text{RINFOG}(7)$ and $\text{RINFOG}(8)$, an estimate for the error in the solution in $\text{RINFOG}(9)$, and condition numbers for the matrix in $\text{RINFOG}(10)$ and $\text{RINFOG}(11)$. See also Section 2.4. Note that if performance is important, $\text{ICNTL}(11)$ should be left set to 0 . Finally, note that, in the current version, if $\text{NRHS} > 1$ or if $\text{ICNTL}(21)=1$ (solution vector kept distributed) then error analysis is not performed and $\text{ICNTL}(11)$ is treated as 0 .

$\text{ICNTL}(12)$ is meaningful only on general symmetric matrices ($\text{SYM} = 2$) and its default value is 0 (automatic choice). For unsymmetric matrices ($\text{SYM}=0$) or symmetric definite positive matrices ($\text{SYM}=1$) all values of $\text{ICNTL}(12)$ are treated as 1 (nothing done). It is only accessed by the host and only during the analysis phase. It defines the ordering strategy (see [18] for more details) and is used, in conjunction with $\text{ICNTL}(6)$, to add constraints to the ordering algorithm. ($\text{ICNTL}(7)$ option). Possible values of $\text{ICNTL}(12)$ are :

- 0 : automatic choice
- 1 : usual ordering (nothing done)
- 2 : ordering on the compressed graph associated with the matrix.
- 3 : constrained ordering, only available with AMF ($\text{ICNTL}(7)=2$).

Other values are treated as 0 . $\text{ICNTL}(12)$, $\text{ICNTL}(6)$, $\text{ICNTL}(7)$ values are strongly related. Therefore, as for $\text{ICNTL}(6)$, if the matrix is in elemental format ($\text{ICNTL}(5)=1$), or the ordering is provided by the user ($\text{ICNTL}(7)=1$), or the Schur option ($\text{ICNTL}(19) \neq 0$) is required, or the matrix is initially distributed ($\text{ICNTL}(18) \neq 0$) then $\text{ICNTL}(12)$ is treated as one.

If MUMPS detects some incompatibility between control parameters then it uses the following rules to automatically reset the control parameters. Firstly $\text{ICNTL}(12)$ has a lower priority than $\text{ICNTL}(7)$ so that if $\text{ICNTL}(12) = 3$ and the ordering required is not AMF then $\text{ICNTL}(12)$ is internally treated as 2 . Secondly $\text{ICNTL}(12)$ has a higher priority than $\text{ICNTL}(6)$ and $\text{ICNTL}(8)$. Thus if $\text{ICNTL}(12) = 2$ and $\text{ICNTL}(6)$ was not active ($\text{ICNTL}(6)=0$) then $\text{ICNTL}(6)$ is automatically reset (treated as $\text{ICNTL}(6)=7$). Furthermore, if $\text{ICNTL}(12) = 3$ then $\text{ICNTL}(6)$ is automatically set to 5 and $\text{ICNTL}(8)$ is set to -2 .

On output from the analysis phase, $\text{INFOG}(24)$ holds the value of $\text{ICNTL}(12)$ that was effectively used. Note that $\text{INFOG}(7)$ and $\text{INFOG}(23)$ hold the values of $\text{ICNTL}(7)$ and $\text{ICNTL}(6)$ (respectively) that were effectively used.

$\text{ICNTL}(13)$ has default value 0 and is only accessed by the host during the analysis phase. If $\text{ICNTL}(13) \leq 0$, ScaLAPACK will be used for the root frontal matrix (last Schur complement to be factored)

if its size is larger than a machine-dependent minimum size. Otherwise (ICNTL(13) > 0), ScaLAPACK will not be used and the root node will be treated sequentially. Processing the root sequentially can be useful when the user is interested in the inertia of the matrix (see INFO(12) and INFOG(12)), or when the user wants to detect null pivots (see ICNTL(24)).

This parameter also controls splitting of the root frontal matrix. If the number of working processors is strictly larger than ICNTL(13) with ICNTL(13)>0 (ScaLAPACK off), then splitting of the root node is performed, in order to automatically recover part of the parallelism lost because the root node was processed sequentially. Finally, setting ICNTL(13) to -1 will force splitting of the root node in all cases (even sequentially), while values strictly smaller than -1 will be treated as 0.

Note that, although ICNTL(13) controls the efficiency of the factorization and solve phases, preprocessing work is performed during analysis and this option must be set on entry to the analysis phase.

ICNTL(14) is accessed by the host both during the analysis and the factorization phases. It corresponds to the percentage increase in the estimated working space. When significant extra fill-in is caused by numerical pivoting, larger values of ICNTL(14) may help use the real/complex working space more efficiently. Except in special cases, the default value is 20 %.

ICNTL(15-17) Not used in current version.

ICNTL(18) has default value 0 and is only accessed by the host during the analysis phase, if the matrix format is assembled (ICNTL(5) = 0). ICNTL(18) defines the strategy for the distributed input matrix. Possible values are:

- 0: the input matrix is centralized on the host. This is the default, see Section 4.5.
- 1: the user provides the structure of the matrix on the host at analysis, MUMPS returns a mapping and the user should then provide the matrix distributed according to the mapping on entry to the numerical factorization phase.
- 2: the user provides the structure of the matrix on the host at analysis, and the distributed matrix on all slave processors at factorization. Any distribution is allowed.
- 3: user directly provides the distributed matrix input both for analysis and factorization.

For options 1, 2, 3, see Section 4.7 for more details on the input/output parameters to MUMPS. For flexibility, options 2 or 3 are recommended.

ICNTL(19) has default value 0 and is only accessed by the host during the analysis phase. If ICNTL(19)=1, then the Schur complement matrix will be returned to the user on the host after the factorization phase. If ICNTL(19)=2 or 3, then the Schur will be returned to the user on the slave processors in the form of a 2D block cyclic distributed matrix (ScaLAPACK style). Values not equal to 1, 2 or 3 are treated as 0. If ICNTL(19) equals 1, 2, or 3, the user must set on entry to the analysis phase, on the host node:

- the integer variable SIZE_SCHUR to the size of the Schur matrix,
- the integer array pointer LISTVAR_SCHUR to the list of indices of the Schur matrix.

For a distributed Schur complement (ICNTL(19)=2 or 3), the integer variables NPROW, NPCOL, MBLOCK, NBLOCK may also be defined on the host before the analysis phase (default values will otherwise be provided). Furthermore, workspace should be allocated by the user before the factorization phase in order for MUMPS to store the Schur complement (see SCHUR, SCHUR_MLOC, SCHUR_NLOC, and SCHUR_LLD in Section 4.10).

Note that the partial factorization of the interior variables can then be exploited to perform a solve phase (transposed matrix or not, see ICNTL(9)). Note that the right-hand side (RHS) provided on input must still be of size N (or $N \times \text{NRHS}$ in case of multiple right-hand sides) even if only the N-SIZE_SCHUR indices will be considered and if only N-SIZE_SCHUR indices of the solution will be relevant to the user.

Finally note that since the Schur complement can be viewed as a partial factorization of the global matrix (with partial ordering of the variables provided by the user) the following options of MUMPS are incompatible with the Schur option: maximum transversal, scaling, iterative refinement, error analysis. Note that if the ordering is given (ICNTL(7)=1) then the following property should hold: $\text{PERM_IN}(\text{LISTVAR_SCHUR}(i)) = \text{N-SIZE_SCHUR} + i$, for $i=1, \text{SIZE_SCHUR}$.

ICNTL(20) has default value 0 and is only accessed by the host during the solve phase. If ICNTL(20)=0, the right-hand side must be given in dense form in the structure component RHS. If ICNTL(20)=1, then the right-hand side must be given in sparse form using the structure components IRHS_SPARSE, RHS_SPARSE, IRHS_PTR and NZ_RHS. Values different from 0 and 1 are treated as 0. (See Section 4.12).

ICNTL(21) has default value 0 and is only accessed by the host during the solve phase. If ICNTL(21)=0, the solution vector will be assembled and stored in the structure component RHS, that must have been allocated earlier by the user. If ICNTL(21)=1, the solution vector is kept distributed at the end of the solve phase, and will be available on each slave processor in the structure components ISOL_loc and SOL_loc. ISOL_loc and SOL_loc must then have been allocated by the user and must be of size at least INFO(23), where INFO(23) has been returned by MUMPS at the end of the factorization phase. Values of ICNTL(21) different from 0 and 1 are currently treated as 0.

Note that if the solution is kept distributed, error analysis and iterative refinement (controlled by ICNTL(10) and ICNTL(11)) are not applied.

ICNTL(22-23) are not used in current version

ICNTL(24) has default value 0 and controls the detection of “null pivot rows”. Null pivot rows are modified to enable the solution phase to provide one solution among the possible solutions of the numerically deficient matrix. Note that the list of row indices corresponding to null pivots is returned on the host in PIVNUL_LIST(1:INFOG(28)). The solution phase (JOB=3) can then be used to either provide a “regular” solution (in the sense that it is a possible solution of the complete system when the right-hand-side belongs to the span of the original matrix) or to compute the associated vectors of the null-space basis (see ICNTL(25)). Combined with the option for null space computation (see ICNTL(25)) during the solution phase, the null space associated to those null pivot rows can then be computed. Possible values of ICNTL(24) are:

- 0 Nothing done. A null pivot will result in error INFO(1)=-10.
- 1 Null pivot row detection; CNTL(3) is used to compute the threshold to decide that a pivot row is “null”. The parameter CNTL(5) then defines the fixation that will be used to enable the solution phase to provide a possible solution to the original system.

Note that when ScaLAPACK is applied on the root node (see ICNTL(13)), then exact null pivots on the root will stop the factorization (INFO(1)=-10) while tiny pivots on the root node will still be factored. Setting ICNTL(13) to a non-zero value will help with the correct detection of null pivots but degrade performance.

ICNTL(25) has default value 0 and is only accessed during the solution phase. It allows the computation of a null space basis, which is meaningful only if a zero-pivot detection option was requested (ICNTL(24) $\neq$ 0) during the factorization and if the matrix was found to be deficient (INFOG(28) > 0); Possible values of ICNTL(25) are:

- 0 A normal solution step is performed. If the matrix was found singular during factorization then one possible solution is returned
- i with $1 \leq i \leq \text{INFOG}(28)$. The i -th vector of the null space basis is computed.
- -1. The complete null space basis is computed.
- Other values result in an error.

Note that when vectors from the null space are requested, both centralized and distributed solutions options can be used. In both cases space to store the null space vectors must be allocated by the user and provided to MUMPS. If the solution is centralized (ICNTL(21)=0), then the null space vectors are returned to the user in the array RHS, allocated by the user on the host. If the solution is distributed (ICNTL(21)=1), then the null space vectors are returned in the array SOL_LOC. In both cases, note that the number of columns of RHS or SOL_LOC must be equal to the number of vectors requested, do that NRHS is equal to:

- 1 if $1 \leq \text{ICNTL}(25) \leq \text{INFOG}(28)$;
- INFOG(28) if ICNTL(25)=-1.

Finally, note that iterative refinement, error analysis, and the option to solve the transpose system (ICNTL(9)) are ignored when the solution step is used to return vectors from the null space (ICNTL(25) $\neq$ 0).

ICNTL(26) has default value 0 and is only accessed by the host during the solution phase. It is only significant if combined with the Schur option (ICNTL(19) $\neq$ 0, see above). It can be used to condense/reduce (ICNTL(26)=1) the right-hand side on the Schur variables, or to expand (ICNTL(26)=2) the Schur local solution on the complete solution (see Section 2.6).

If ICNTL(26) $\neq$ 0, then the user should provide workspace in the pointer array REDRHS, as well as a leading dimension LREDRHS (see Section 4.10).

If ICNTL(26)=1 then only a forward substitution is performed. The solution corresponding to the ‘internal’ (non-Schur) variables is returned together with the reduced/condensed right-hand-side. The reduced right-hand side is made available on the host in REDRHS.

If ICNTL(26)=2 then REDRHS is considered to be the solution corresponding to the Schur variables. The backward substitution is then performed with the given right-hand side to compute the solution associated with the ‘internal’ variables. Note that the solution corresponding to the Schur variables is also made available in the main solution vector/matrix.

Values different from 1 and 2 are treated as 0. Note that if no Schur complement was computed, ICNTL(26) = 1 or 2 results in an error. Finally, if ICNTL(26) = 1 or 2, then error analysis and iterative refinements are disabled.

ICNTL(27-40) are not used in the current version.

mumps_par%**CNTL** is a **real** (also **real** in the complex version) array of dimension 5.

CNTL(1) is the relative threshold for numerical pivoting. It forms a trade-off between preserving sparsity and ensuring numerical stability during the factorization. In general, a larger value of CNTL(1) increases fill-in but leads to a more accurate factorization. If CNTL(1) is nonzero, numerical pivoting will be performed. If CNTL(1) is zero, no such pivoting will be performed and the subroutine will fail if a zero pivot is encountered. If the matrix is diagonally dominant, then setting CNTL(1) to zero will decrease the factorization time while still providing a stable decomposition. If the code is called for unsymmetric or general symmetric matrices, CNTL(1) has default value 0.01. For symmetric positive definite matrices numerical pivoting is suppressed and the default value is 0.0. Values less than 0.0 are treated as 0.0. In the unsymmetric case (respectively symmetric case), values greater than 1.0 (respectively 0.5) are treated as 1.0 (respectively 0.5).

CNTL(2) is the stopping criterion for iterative refinement and is only accessed by the host during the solve phase. Let $Berr = \max_i \frac{|r|_i}{(|A| \cdot |x| + |b|)_i}$ [9]. Iterative refinement will stop when either the required accuracy is reached ($Berr < \text{CNTL}(2)$) or the convergence rate is too slow ($Berr$ does not decrease by at least a factor of 5). Default value is $\sqrt{\epsilon}$.

CNTL(3) is only used combined with null pivot detection (ICNTL(24) = 1) and is not used otherwise. CNTL(3) has default value -1.0 and is only accessed by the host during the numerical factorization phase. A pivot is considered to be null if its row/column is smaller than $thres \times \|A\|$, where $thres = \text{CNTL}(3)$ if $\text{CNTL}(3) > 0$. If $\text{CNTL}(3) \leq 0$, then $thres = \epsilon \times 10^{-5} \times \|A\|$, where ϵ holds the machine precision, is used.

CNTL(4) determines the threshold for static pivoting. It is only accessed by the host, and must be set either before the factorization phase, or before the analysis phase. It has default value -1.0. If $\text{CNTL}(4) < 0.0$ static pivoting is not activated. If $\text{CNTL}(4) > 0.0$ static pivoting is activated and the magnitude of small pivots smaller than CNTL(4) will be set to CNTL(4). If $\text{CNTL}(4) = 0.0$ static pivoting is activated and the threshold value used is determined automatically.

CNTL(5) is the fixation for null pivots and is effective only when null pivot detection is active (ICNTL(24) = 1). CNTL(5) has default value 0.0 and is only accessed by the host during the numerical factorization phase. If $\text{CNTL}(5) > 0$ the detected null pivot is set to $\text{CNTL}(5) \times \|A\|$. Furthermore, the sign of the pivot is preserved in the modified diagonal entry. If $\text{CNTL}(5) \leq 0$, then the pivot row (except the pivot) is set to zero and the pivot is set to one. In symmetric case, the pivot column (except the pivot) is also set to 0.

CNTL(6-15) are not used in the current version.

6 Information parameters

The parameters described in this section are returned by MUMPS and hold information that may be of interest to the user. Some of the information is local to each processor and some only on the host. If an error is detected (see Section 7), the information may be incomplete.

6.1 Information local to each processor

The arrays `mumps_par%RINFO` and `mumps_par%INFO` are local to each process.

`mumps_par%RINFO` is a double precision array of dimension 20. It contains the following local information on the execution of MUMPS:

- RINFO(1) - after analysis: The estimated number of floating-point operations on the processor for the elimination process.
- RINFO(2) - after factorization: The number of floating-point operations on the processor for the assembly process.
- RINFO(3) - after factorization: The number of floating-point operations on the processor for the elimination process.
- RINFO(4) - RINFO(20) are not used in the current version.

`mumps_par%INFO` is an integer array of dimension 40. It contains the following local information on the execution of MUMPS:

- INFO(1) is 0 if the call to MUMPS was successful, negative if an error occurred (see Section 7), or positive if a warning is returned.
- INFO(2) holds additional information about the error or the warning. If INFO(1) = -1, INFO(2) is the processor number (in communicator `mumps_par%COMM`) on which the error was detected.
- INFO(3) - after analysis: Estimated real space needed on the processor for the factors.
- INFO(4) - after analysis: Estimated integer space needed on the processor for factors.
- INFO(5) - after analysis: Estimated maximum front size on the processor.
- INFO(6) - after analysis: Number of nodes in the complete tree. The same value is returned on all processors.
- INFO(7) - after analysis: Minimum value of MAXIS estimated by the analysis phase to run the numerical factorization successfully.
- INFO(8) - after analysis: Minimum value of MAXS estimated by the analysis phase to run the numerical factorization successfully.
- INFO(9) - after factorization: Size of the real space used on the processor to store the factors.
- INFO(10) - after factorization: Size of the integer space used on the processor to store the factors.
- INFO(11) - after factorization: Order of the largest frontal matrix processed on the processor.
- INFO(12) - after factorization: Number of off-diagonal pivots selected on the processor if SYM=0 or number of negative pivots on the processor if SYM=1 or 2. If ICNTL(13)=0 (the default), this excludes pivots from the parallel root node treated by ScaLAPACK. (This means that the user should set ICNTL(13)=1 or use a single processor in order to get the exact number of off-diagonal or negative pivots rather than a lower bound.) Note that if SYM=1 or 2, INFO(12) will be 0 for complex symmetric matrices.
- INFO(13) - after factorization: The number of uneliminated variables, corresponding to delayed pivots, sent to the father. If a delayed pivot is subsequently passed to the father of the father, it is counted a second time.
- INFO(14) - after factorization: Number of memory compresses on the processor.
- INFO(15) - after analysis: estimated total size (in millions of bytes) of all MUMPS internal data for running numerical factorization.

INFO(16) - after factorization: total size (in millions of bytes) of all MUMPS internal data used during numerical factorization.

INFO(18) - after factorization: local number of null pivots detected when ICNTL(24)=1.

INFO(19) - INFO(22) are not used in the current version. Are set to zero.

INFO(23) - after factorization: total number of pivots eliminated on the processor. In the case of a distributed solution (see ICNTL(21)), this should be used by the user to allocate solution vectors ISOL_loc and SOL_loc of appropriate dimensions (ISOL_LOC of size INFO(23), SOL_LOC of size LSOL_LOC $\times$ NRHS where LSOL_LOC $\geq$ INFO(23)) on that processor, between the factorization and solve steps.

INFO(24) - INFO(40) are not used in the current version.

6.2 Information available on all processors

The arrays mumps_par%RINFOG and mumps_par%INFOG :

mumps_par%**RINFOG** is a double precision array of dimension 20. It contains the following global information on the execution of MUMPS:

RINFOG(1) - after analysis: The estimated number of floating-point operations (on all processors) for the elimination process.

RINFOG(2) - after factorization: The total number of floating-point operations (on all processors) for the assembly process.

RINFOG(3) - after factorization: The total number of floating-point operations (on all processors) for the elimination process.

RINFOG(4) to RINFOG(11) - after solve with error analysis: Only returned if ICNTL(11) $\neq$ 0. See description of ICNTL(11).

RINFOG(12) - RINFOG(20) are not used in the current version.

mumps_par%**INFOG** is an integer array of dimension 40. It contains the following global information on the execution of MUMPS:

INFOG(1) is 0 if the call to MUMPS was successful, negative if an error occurred (see Section 7), or positive if a warning is returned.

INFOG(2) holds additional information about the error or the warning.

The difference between INFOG(1:2) and INFO(1:2) is that INFOG(1:2) is the same on all processors. It has the value of INFO(1:2) of the processor which returned with the most negative INFO(1) value. For example, if processor p returns with INFO(1)=-13, and INFO(2)=10000, then all other processors will return with INFOG(1)=-13 and INFOG(2)=10000, but still INFO(1)=-1 and INFO(2)= p .

INFOG(3) - after analysis: Total estimated real/complex workspace for factors on all processors. If negative, then the absolute value corresponds to *millions* of real/complex entries in the workspace allocated for the factors. (Note that for symmetric matrices and in parallel, this can be larger than INFOG(20).)

INFOG(4) - after analysis: Total estimated integer workspace for factors on all processors.

INFOG(5) - after analysis: Estimated maximum front size in the complete tree.

INFOG(6) - after analysis: Number of nodes in the complete tree.

INFOG(7) - after analysis: ordering option effectively used (see ICNTL(7)).

INFOG(8) - after analysis: structural symmetry in percent (100 : symmetric, 0 : fully unsymmetric) of the (permuted) matrix. (-1 indicates that the structural symmetry was not computed which will be the case if the input matrix is in elemental form.)

INFOG(9) - after factorization: Total real/complex workspace to store the LU factors. If negative, then the absolute value corresponds to the size in *millions* of real/complex entries used to store the factors.

INFOG(10) - after factorization: Total integer space to store the LU factors.

INFOG(11) - after factorization: Order of largest frontal matrix.

INFOG(12) - after factorization: Total number of off-diagonal pivots if SYM=0 or total number of negative pivots (real arithmetic) if SYM=1 or 2. If ICNTL(13)=0 (the default) this excludes pivots from the parallel root node treated by ScaLAPACK. (This means that the user should set ICNTL(13) to a positive value, say 1, or use a single processor in order to get the exact number of off-diagonal or negative pivots rather than a lower bound.) Note that if SYM=1 or 2, INFOG(12) will be 0 for complex symmetric matrices.

INFOG(13) - after factorization: Total number of delayed pivots.

INFOG(14) - after factorization: Total number of memory compresses.

INFOG(15) - after solution: Number of steps of iterative refinement.

INFOG(16) - after analysis: Estimated size (in million of bytes) of all MUMPS internal data for running factorization (value on the most memory consuming processor).

INFOG(17) - after analysis: Estimated size (in millions of bytes) of all MUMPS internal data for running factorization (sum over all processors).

INFOG(18) - after factorization: Size in millions of bytes of all MUMPS internal data allocated during factorization: value on the most memory consuming processor.

INFOG(19) - after factorization: Size in millions of bytes of all MUMPS internal data allocated during factorization: sum over all processors.

INFOG(20) - after analysis: Estimated number of entries in the factors. If negative the absolute value corresponds to *millions* of entries in the factors.

INFOG(21) - after factorization: Size in millions of bytes of memory effectively used during factorization: value on the most memory consuming processor.

INFOG(22) - after factorization: Size in millions of bytes of memory effectively used during factorization: sum over all processors.

INFOG(23) - After analysis: value of ICNTL(6) effectively used.

INFOG(24) - After analysis: value of ICNTL(12) effectively used.

INFOG(25) - After factorization : number of tiny pivots (number of pivots modified by static pivoting)

INFOG(26-27) are not used in the current version.

INFOG(28) - After factorization : number of null pivots encountered. See CNTL(3) for the definition of a null pivot.

INFOG(29) - INFOG(40) are not used in the current version.

7 Error diagnostics

MUMPS uses the following mechanism to process errors that may occur during the parallel execution of the code. If, during a call to MUMPS, an error occurs on a processor, this processor informs all the other processors before they return from the call. In parts of the code where messages are sent asynchronously (for example the factorization and solve phases), the processor on which the error occurs sends a message to the other processors with a specific error tag. On the other hand, if the error occurs in a subroutine that does not use asynchronous communication, the processor propagates the error to the other processors.

On successful completion, a call to MUMPS will exit with the parameter `mumps_par%INFOG(1)` set to zero. A negative value for `mumps_par%INFOG(1)` indicates that an error has been detected on one of the processors. For example, if processor *s* returns with `INFO(1) = -8` and `INFO(2)=1000`, then processor *s* ran out of integer workspace during the factorization and the size of the workspace `MAXIS` should be increased by 1000 at least. The other processors are informed about this error and return with `INFO(1) = -1` (i.e., an error occurred on another processor) and `INFO(2)=s` (i.e., the error occurred on processor *s*). Processors that detected a local error, do not overwrite `INFO(1)`, i.e., only processors that did not produce an error will set `INFO(1)` to `-1` and `INFO(2)` to the processor having the most negative error code.

The behaviour is slightly different for INFO(1) and INFO(2): in the previous example, all processors would return with INFO(1) = -8 and INFO(2)=1000.

The possible error codes returned in INFO(1) (and INFO(2)) have the following meaning:

- 1 An error occurred on processor INFO(2).
- 2 NZ is out of range. INFO(2)=NZ.
- 3 MUMPS was called with an invalid value for JOB. This may happen for example if the analysis (JOB=1) was not performed before the factorization (JOB=2), or the factorization was not performed before the solve (JOB=3), or the initialization phase (JOB=-1) was performed a second time on an instance not freed (JOB=-2). See description of JOB in Section 3. This error also occurs if JOB does not contain the same value on all processes on entry to MUMPS.
- 4 Error in user-provided permutation array PERM_LIN in position INFO(2). This error occurs on the host only.
- 5 Problem of REAL workspace allocation of size INFO(2) during analysis.
- 6 Matrix is singular in structure.
- 7 Problem of INTEGER workspace allocation of size INFO(2) during analysis.
- 8 MAXIS too small for factorization. This may happen, for example, if numerical pivoting leads to significantly more fill-in than was predicted by the analysis. The user should increase the value of ICNTL(14) or the value of MAXIS before recalling the factorization (JOB=2).
- 9 MAXS too small for factorization. The user should increase the value of ICNTL(14) or MAXS before recalling the factorization (JOB=2).
- 10 Numerically singular matrix.
- 11 MAXS too small for solution. See error INFO(1) = -9.
- 12 MAXS too small for iterative refinement. See error INFO(1) = -9.
- 13 Error in a Fortran ALLOCATE statement. INFO(2) contains the size that the package requested.
- 14 MAXIS too small for solution. See error INFO(1) = -8.
- 15 MAXIS too small for iterative refinement and/or error analysis. See error INFO(1) = -8.
- 16 N is out of range. INFO(2)=N.
- 17 The internal send buffer that was allocated dynamically by MUMPS on the processor is too small. The user should increase the value of ICNTL(14) before calling MUMPS again.
- 18 MAXIS too small to process root node. See error INFO(1) = -8.
- 19 MAXS too small to process root node. See error INFO(1) = -9.
- 20 The internal reception buffer that was allocated dynamically by MUMPS is too small. INFO(2) holds the minimum size of the reception buffer required (in bytes). The user should increase the value of ICNTL(14) before calling MUMPS again.
- 21 Value of PAR=0 is not allowed because only one processor is available; INFO(2) is set to the number of processors, 1. Running MUMPS in host-node mode (the host is not a slave processor itself) requires at least two processors. The user should either set PAR to 1 or increase the number of processors.
- 22 A pointer array is provided by the user that is either
 - not associated, or
 - has insufficient size, or
 - is associated and should not be associated (for example, RHS on non-host processors).
 INFO(2) points to the pointer array having the wrong format in the table below:

INFO(2)	array
1	IRN or ELTPTR
2	JCN or ELTVAR
3	PERM_IN
4	A or A.ELT
5	ROWSCA
6	COLSCA
7	RHS
8	LISTVAR_SCHUR
9	SCHUR
10	RHS_SPARSE
11	IRHS_SPARSE
12	IRHS_PTR
13	ISOL_LOC
14	SOL_LOC
15	REDRHS

- 23 MPI was not initialized by the user prior to a call to MUMPS with JOB = -1.
 - 24 NELT is out of range. INFO(2)=NELT.
 - 25 A problem has occurred in the initialization of the BLACS. This may be because you are using a vendor's BLACS. Try using a BLACS version from netlib instead.
 - 26 LRHS is out of range. INFO(2) = LRHS.
 - 27 NZ_RHS and IRHS_PTR(NRHS+1) do not match. INFO(2) = IRHS_PTR(NRHS+1).
 - 28 IRHS_PTR(1) is not equal to 1. INFO(2) = IRHS_PTR(1).
 - 29 LSOL_LOC is smaller than KEEP(89). INFO(2)=LSOL_LOC.
 - 30 SCHUR_LLD is out of range. INFO(2) = SCHUR_LLD.
 - 31 A 2D block cyclic Schur complement is required with the option ICNTL(19)=3, but the user has provided a process grid that does not satisfy the constraint MBLOCK=NBLOCK. INFO(2)=MBLOCK-NBLOCK.
 - 32 Incompatible values of NRHS and ICNTL(25). Either ICNTL(25) was set to -1 and NRHS is different from INFOG(28); or ICNTL(25) was set to i , $1 \leq i \leq \text{INFOG}(28)$ and NRHS is different from 1. Value of NRHS is stored in INFO(2).
 - 33 ICNTL(26) was asked during solve phase but Schur complement was not computed during factorization. INFO(2)=ICNTL(26).
 - 34 LREDRHS is out of range. INFO(2)=LRHS.
 - 35 Expansion phase was called (ICNTL(26) = 2) but reduction phase (ICNTL(26)=1) was not called before.
 - 36 Incompatible values of ICNTL(25) and INFOG(28). Value of ICNTL(25) is stored in INFO(2).
- A positive value of INFO(1) is associated with a warning message which will be output on unit ICNTL(2) when ICNTL(4) $\geq$ 2.
- +1 Index (in IRN or JCN) out of range. Action taken by subroutine is to ignore any such entries and continue. INFO(2) is set to the number of faulty entries. Details of the first ten are printed on unit ICNTL(2).
 - +2 During error analysis the max-norm of the computed solution was found to be zero.
 - +4 User data JCN has been corrupted.
 - +8 Warning return from the iterative refinement routine. More than ICNTL(10) iterations are required.
 - + Combinations of the above warnings will correspond to summing the constituent warnings.

8 Calling MUMPS from C

MUMPS is a Fortran 90 library, designed to be used from Fortran 90 rather than C. However a basic C interface is provided that allows users to call MUMPS directly from C programs. Similarly to the Fortran 90 interface, the C interface uses a structure whose components match those in the MUMPS structure for Fortran (Figure 1). Thus the description of the parameters in Sections 4 and 5 applies. Figure 2 shows the C structure [SDCZ]MUMPS_STRUC_C. This structure is defined in the include file [sdcz]mumps_c.h and there is one main routine per available precision with the following prototype:

```
void [sdcz]mumps_c(MUMPS_STRUC_C * idptr);
```

```
typedef struct
{
    int sym, par, job;
    int comm_fortran; /* Fortran communicator */
    int icntl[40];
    real cntl[15];
    int n;
    /* Assembled entry */
    int nz; int *irn; int *jcn; real/complex *a;
    /* Distributed entry */
    int nz_loc; int *irn_loc; int *jcn_loc; real/complex *a_loc;
    /* Element entry */
    int nelt; int *eltptr; int *eltvar; real/complex *a_elt;
    /* Ordering, if given by user */
    int *perm_in;
    /* Scaling (input only in this version) */
    real/complex *colsca; real/complex *rowsca;
    /* RHS, solution, output data and statistics */
    real/complex *rhs, *redrhs, *rhs_sparse, *sol_loc;
    int *irhs_sparse, *irhs_ptr, *isol_loc;
    int nrhs, lrhs, lredrhs, nz_rhs, lsol_loc;
    int info[40], infog[40];
    real rinfo[20], rinfog[20];
    int *sym_perm, *uns_perm;
    int *mapping;
    /* Schur */
    int size_schur; int *listvar_schur; real/complex *schur;
    int nprow, npcol, mblock, nblock, schur_lld, schur_mloc, schur_nloc;
    /* Version number */
    char version_number[80];
    /* Internal parameters */
    int instance_number;
} [SDCZ]MUMPS_STRUC_C;
```

Figure 2: Definition of the C structure [SDCZ]MUMPS_STRUC_C. **real/complex** is used for data that can be either real or complex, **real** for data that stays real (**float** or **double**) in the complex version.

An example of calling MUMPS from C for a complex assembled problem is given in Section 10.3. The following subsections discuss some technical issues that a user should be aware of before using the C interface to MUMPS.

In the following, we suppose that `id` has been declared of type [SDCZ]MUMPS_STRUC_C.

8.1 Array indices

Arrays in C start at index 0 whereas they normally start at 1 in Fortran. Therefore, care must be taken when providing arrays to the C structure. For example, the row indices of the matrix `A`, stored in `IRN(1:NZ)` in the Fortran version should be stored in `irn[0:nz-1]` in the C version. (Note that the contents of `irn` itself is unchanged with values between 1 and `N`.) One solution to deal with this is to define macros:

```
#define ICNTL( i ) icntl[ (i) - 1 ]
#define A( i ) a[ (i) -1 ]
#define IRN( i ) irn[ (i) -1 ]
...
```

and then use the uppercase notation with parenthesis (instead of lowercase/brackets). In that case, the notation `id.IRN(I)`, where `I` is in $\{1, 2, \dots, \text{NZ}\}$ can be used instead of `id.irn[I-1]`; this notation then matches exactly with the description in Sections 4 and 5, where arrays are supposed to start at 1.

This can be slightly more confusing for element matrix input (see Section 4.6), where some arrays are used to index other arrays. For instance, the first value in `eltptr`, `eltptr[0]`, pointing into the list of variables of the first element in `eltvar`, should be equal to 1. Effectively, using the notation above, the list of variables for element $j = 1$ starts at location `ELTVAR(ELTPTR(j)) = ELTVAR(eltptr[j-1]) = eltvar[eltptr[j-1]-1]`.

8.2 Issues related to the C and Fortran communicators

In general, C and Fortran communicators have a different datatype and are not directly compatible. For the C interface, MUMPS requires a Fortran communicator to be provided in `id.comm.fortran`. If, however, this field is initialized to the special value -987654, the Fortran communicator `MPI_COMM_WORLD` is used by default. If you need to call MUMPS based on a smaller number of processors defined by a C subcommunicator, then you should convert your C communicator to a Fortran one. This has not been included in MUMPS because it is dependent on the MPI implementation and thus not portable. For MPI2, and most MPI implementations, you may just do

```
id.comm_fortran = (F_INT) MPI_Comm_c2f(comm_c);
```

(Note that `F_INT` is defined in `[sdcz]mumps.c.h` and normally is an int.) For MPI implementations where the Fortran and the C communicators have the same integer representation

```
id.comm_fortran = (F_INT) comm_c;
```

should work.

For some MPI implementations, check if `id.comm_fortran = MPIR_FromPointer(comm_c)` can be used.

8.3 Fortran I/O

Diagnostic, warning and error messages (controlled by `ICNTL(1:4) / icntl[0..3]`) are based on Fortran file units. Use the value 6 for the Fortran unit 6 which corresponds to `stdout`. For a more general usage with specific file names from C, passing a C file handler is not currently possible. One solution would be to use a Fortran subroutine along the lines of the model below:

```
SUBROUTINE OPENFILE( UNIT, NAME )
INTEGER UNIT
CHARACTER*(*) NAME
OPEN(UNIT, file=NAME)
RETURN
END
```

and have (in the C user code) a statement like

```
openfile_( &mumps_par.ICNTL(1), name, name_length_byval)
```

(or slightly different depending on the C-Fortran calling conventions); something similar could be done to close the file.

8.4 Runtime libraries

The Fortran 90 runtime library corresponding to the compiler used to compile MUMPS is required at the link stage. One way to provide it is to perform the link phase with the Fortran compiler (instead of the C compiler or `ld`).

8.5 Integer, real and complex datatypes in C and Fortran

We assume that the `int`, `float` and `double` types are compatible with the Fortran `INTEGER`, `REAL` and `DOUBLE PRECISION` datatypes. If this were not the case, the files `[dscz]mumps_prec.h` or Makefiles would need to be modified accordingly.

Since not all C compilers define the `complex` datatype (this only appeared in the C99 standard), we define the following, compatible with the Fortran `COMPLEX` and `DOUBLE COMPLEX` types:

```
typedef struct {float r,i;} mumps_complex; for simple precision (cmumps), and
typedef struct {double r,i;} mumps_double_complex; for double precision
(zmumps).
```

Types for complex data from the user program should be compatible with those above.

8.6 Sequential version

The C interface to MUMPS is compatible with the sequential version; see Section 2.9.

9 Scilab and MATLAB interfaces

The main callable functions are

```
id = initmumps;
id = dmumps(id [,mat] );
id = zmumps(id [,mat] );
```

We have designed these interfaces such that their usage is as similar as possible to the existing C and Fortran interfaces to MUMPS, and where only the parameters related to the sequential code are used. The main differences and characteristics are:

- The existence of a function `initmumps` (usage: `id=initmumps`) that builds an initial structure `id` in which `id.JOB` is set to -1 and `id.SYM` is set to 0 (unsymmetric solver by default).
- Only the double precision and double complex versions of MUMPS are interfaced, since they correspond to the arithmetic precisions used in MATLAB/Scilab.
- the sparse matrix A is passed to the interface functions `dmumps` and `zmumps` as a Scilab/MATLAB object (parameters `ICNTL(5)`, `N`, `NZ`, `NELT`, ... are thus irrelevant).
- the right-hand side vector or matrix, possibly sparse, is passed to the interface functions `dmumps` and/or `zmumps` in the argument `id.RHS`, as a Scilab/MATLAB object (parameters `ICNTL(20)`, `NRHS`, `NZRHS`, ... are thus irrelevant).
- The Schur complement matrix, if required, is allocated within the interface and returned as a Scilab/MATLAB dense matrix. Furthermore, the parameters `SIZE_SCHUR` and `ICNTL(19)` need not be set by the user; they are set automatically depending on the availability and size of the list of Schur variables, `id.VAR_SCHUR`.
- We have chosen to use a new variable `id.SOL` to store the solution, instead of overwriting `id.RHS`.

Please refer to the report [19] for a more detailed description of these interfaces. Please also refer to the README file in directories MATLAB or Scilab of the main MUMPS distribution for more information on installation. For example, one important thing to note is that at installation, the user must provide the Fortran 90 runtime libraries corresponding to the compiled MUMPS package. This can be done in the makefile for the MATLAB interface (file `make.inc`) and in the builder for the Scilab interface (file `builder.sce`).

Finally, note that examples of usage of the MATLAB and the Scilab interfaces are provided in directories `MATLAB` and `Scilab/examples`, respectively. In the following, we describe the input and output parameters of the function `[dz]mumps`, that are relevant in the context of this interface to the sequential version of MUMPS.

Input Parameters

- **mat** : sparse matrix which has to be provided as the second argument of `dmumps` if `id.JOB` is strictly larger than 0.
- **id.SYM** : controls the matrix type (symmetric positive definite, symmetric indefinite or unsymmetric) and it has to be initialized by the user before the initialization phase of MUMPS (see `id.JOB`). Its value is set to 0 after the call of `initmumps`.
- **id.JOB** : defines the action that will be realized by MUMPS: initialize, analyze and/or factorize and/or solve and release MUMPS internal C/Fortran data. It has to be set by the user before any call to MUMPS (except after a call to `initmumps`, which sets its value to -1).
- **id.ICNTL** and **id.CNTL** : define control parameters that can be set after the initialization call (`id.JOB = -1`). See Section “Control parameters” for more details. If the user does not modify an entry in `id.ICNTL` then MUMPS uses the default parameter. For example, if the user wants to use the AMD ordering, he/she should set `id.ICNTL(7) = 0`. Note that the following parameters are inhibited because they are automatically set within the interface: `id.ICNTL(19)` which controls the Schur complement option and `id.ICNTL(20)` which controls the format of the right-hand side. Note that parameters `id.ICNTL(1:4)` may not work properly depending on your compiler and your environment. In case of problem, we recommend to switch printing off by setting `id.ICNTL(1:4)=-1`.
- **id.PERM_IN** : corresponds to the given ordering option (see Section “Input and output parameters” for more details). Note that this permutation is only accessed if the parameter `id.ICNTL(7)` is set to 1.
- **id.COLSCA** and **id.ROWSCA** : are optional scaling arrays (see Section “Input and output parameters” for more details)
- **id.RHS** : defines the right-hand side. The parameter `id.ICNTL(20)` related to its format (sparse or dense) is automatically set within the interface. Note that `id.RHS` is not modified (as in MUMPS), the solution is returned in `id.SOL`.
- **id.VAR_SCHUR** : corresponds to the list of variables that appear in the Schur complement matrix (see Section “Input and output parameters” for more details).
- **id.REDRHS** (input parameter only if `id.VAR_SCHUR` was provided during the factorization and if `ICNTL(26)=2` on entry to the solve phase): partial solution on the variables corresponding to the Schur complement. It is provided by the user and normally results from both the Schur complement and the reduced right-hand side that were returned by MUMPS in a previous call. When `ICNTL(26)=2`, MUMPS uses this information to build the solution `id.SOL` on the complete problem. See Section “Schur complement” for more details.

Output Parameters

- **id.SCHUR** : if `id.VAR_SCHUR` is provided of size `SIZE_SCHUR`, then `id.SCHUR` corresponds to a dense array of size `(SIZE_SCHUR,SIZE_SCHUR)` that holds the Schur complement matrix (see Section “Input and output parameters” for more details). The user does not have to initialize it.
- **id.REDRHS** (output parameter only if `ICNTL(26)=1` and `id.VAR_SCHUR` was defined): Reduced right-hand side (or condensed right-hand side on the variables associated to the Schur complement). It is computed by MUMPS during the solve stage if `ICNTL(26)=1`. It can then be used outside MUMPS, together with the Schur complement, to build a solution on the interface. See Section “Schur complement” for more details.
- **id.INFOG** and **id.RINFOG** : information parameters (see Section “Information parameters”).

- **id.SYM_PERM** : corresponds to a symmetric permutation of the variables (see discussion regarding ICNTL(7) in Section “Control parameters”). This permutation is computed during the analysis and is followed by the numerical factorization except when numerical pivoting occurs.
- **id.UNS_PERM** : column permutation (if any) on exit from the analysis phase of MUMPS (see discussion regarding ICNTL(6) in Section “Control parameters”).
- **id.SOL** : dense vector or matrix containing the solution after MUMPS solution phase.

Internal Parameters

- **id.INST**: (MUMPS reserved component) MUMPS internal parameter.
- **id.TYPE**: (MUMPS reserved component) defines the arithmetic (complex or double precision).

10 Examples of use of MUMPS

10.1 An assembled problem

An example program illustrating a possible use of MUMPS on assembled DOUBLE PRECISION problems is given Figure 3. Two files must be included in the program: `mpif.h` for MPI and `mumps_struct.h` for MUMPS. The file `mumps_root.h` must also be available because it is included in `mumps_struct.h`. The initialization and termination of MPI are performed in the user program via the calls to `MPI_INIT` and `MPI_FINALIZE`.

The MUMPS package is initialized by calling MUMPS with `JOB = -1`, the problem is read in by the host (in the components N, NZ, IRN, JCN, A, and RHS), and the solution is computed in RHS with a call on all processors to MUMPS with `JOB=6`. Finally, a call to MUMPS with `JOB = -2` is performed to deallocate the data structures used by the instance of the package.

Thus for the assembled 5×5 matrix and right-hand side

$$\begin{pmatrix} 2 & 3 & 4 & & \\ 3 & & -3 & 6 & \\ & -1 & 1 & 2 & \\ & & 2 & & \\ & 4 & & & 1 \end{pmatrix}, \quad \begin{pmatrix} 20 \\ 24 \\ 9 \\ 6 \\ 13 \end{pmatrix}$$

we could have as input

```

5          : N
12         : NZ
1 2 3.0
2 3 -3.0
4 3 2.0
5 5 1.0
2 1 3.0
1 1 2.0
5 2 4.0
3 4 2.0
2 5 6.0
3 2 -1.0
1 3 4.0
3 3 1.0    : A
20.0
24.0
9.0
6.0
13.0      : RHS
```

and we obtain the solution $\text{RHS}(i) = i$, $i = 1, \dots, 5$.

```

PROGRAM MUMPS_EXAMPLE
INCLUDE 'mpif.h'
INCLUDE 'dmumps_struct.h'
TYPE (DMUMPS_STRUC) id
INTEGER IERR, I
CALL MPI_INIT(IERR)
C Define a communicator for the package
id%COMM = MPI_COMM_WORLD
C Ask for unsymmetric code
id%SYM = 0
C Host working
id%PAR = 1
C Initialize an instance of the package
id%JOB = -1
CALL DMUMPS(id)
C Define problem on the host (processor 0)
IF ( id%MYID .eq. 0 ) THEN
  READ(5,*) id%N
  READ(5,*) id%NZ
  ALLOCATE( id%IRN ( id%NZ ) )
  ALLOCATE( id%JCN ( id%NZ ) )
  ALLOCATE( id%A( id%NZ ) )
  ALLOCATE( id%RHS ( id%N ) )
  READ(5,*) ( id%IRN(I) ,I=1, id%NZ )
  READ(5,*) ( id%JCN(I) ,I=1, id%NZ )
  READ(5,*) ( id%A(I),I=1, id%NZ )
  READ(5,*) ( id%RHS(I) ,I=1, id%N )
END IF
C Call package for solution
id%JOB = 6
CALL DMUMPS(id)
C Solution has been assembled on the host
IF ( id%MYID .eq. 0 ) THEN
  WRITE( 6, * ) ' Solution is ',(id%RHS(I),I=1,id%N)
END IF
C Deallocate user data
IF ( id%MYID .eq. 0 )THEN
  DEALLOCATE( id%IRN )
  DEALLOCATE( id%JCN )
  DEALLOCATE( id%A )
  DEALLOCATE( id%RHS )
END IF
C Destroy the instance (deallocate internal data structures)
id%JOB = -2
CALL DMUMPS(id)
CALL MPI_FINALIZE(IERR)
STOP
END

```

Figure 3: Example program using MUMPS on an assembled DOUBLE PRECISION problem

10.2 An elemental problem

An example of a driver to use MUMPS for element DOUBLE PRECISION problems is given in Figure 4. The calling sequence is similar to that for the assembled problem in Section 10.1 but now the host reads the problem in components N, NELT, ELTPTR, ELTVAR, A.ELT, and RHS. Note that for elemental problems ICNTL(5) must be set to 1 and that elemental matrices always have a symmetric structure. For the two-element matrix and right-hand side

$$\begin{matrix} 1 \\ 2 \\ 3 \end{matrix} \begin{pmatrix} -1 & 2 & 3 \\ 2 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}, \quad \begin{matrix} 3 \\ 4 \\ 5 \end{matrix} \begin{pmatrix} 2 & -1 & 3 \\ 1 & 2 & -1 \\ 3 & 2 & 1 \end{pmatrix}, \quad \begin{pmatrix} 12 \\ 7 \\ 23 \\ 6 \\ 22 \end{pmatrix}$$

we could have as input

```
5
2
6
18
1 4 7
1 2 3 3 4 5
-1.0 2.0 1.0 2.0 1.0 1.0 3.0 1.0 1.0 2.0 1.0 3.0 -1.0 2.0 2.0 3.0 -1.0 1.0
12.0 7.0 23.0 6.0 22.0
```

and we obtain the solution $\text{RHS}(i) = i$, $i = 1, \dots, 5$.

10.3 An example of calling MUMPS from C

An example of a driver to use MUMPS from C is given in Figure 5.

```

PROGRAM MUMPS_EXAMPLE
  INCLUDE 'mpif.h'
  INCLUDE 'dmumps_struct.h'
  TYPE (DMUMPS_STRUC) id
  INTEGER IERR, LEFTVAR, NA_ELT
  CALL MPI_INIT(IERR)
C Define a communicator for the package
  id%COMM = MPI_COMM_WORLD
C Ask for unsymmetric code
  id%SYM = 0
C Host working
  id%PAR = 1
C Initialize an instance of the package
  id%JOB = -1
  CALL DMUMPS(id)
C Define the problem on the host (processor 0)
  IF ( id%MYID .eq. 0 ) THEN
    READ(5,*) id%N
    READ(5,*) id%NELT
    READ(5,*) LEFTVAR
    READ(5,*) NA_ELT
    ALLOCATE( id%ELTPTR ( id%NELT+1 ) )
    ALLOCATE( id%ELTVAR ( LEFTVAR ) )
    ALLOCATE( id%A_ELT( NA_ELT ) )
    ALLOCATE( id%RHS ( id%N ) )
    READ(5,*) ( id%ELTPTR(I) ,I=1, id%NELT+1 )
    READ(5,*) ( id%ELTVAR(I) ,I=1, LEFTVAR )
    READ(5,*) ( id%A_ELT(I),I=1, NA_ELT )
    READ(5,*) ( id%RHS(I) ,I=1, id%N )
  END IF
C Specify element entry
  id%ICNTL(5) = 1
C Call package for solution
  id%JOB = 6
  CALL DMUMPS(id)
C Solution has been assembled on the host
  IF ( id%MYID .eq. 0 ) THEN
    WRITE( 6, * ) ' Solution is ',(id%RHS(I),I=1,id%N)
C Deallocate user data
    DEALLOCATE( id%ELTPTR )
    DEALLOCATE( id%ELTVAR )
    DEALLOCATE( id%A_ELT )
    DEALLOCATE( id%RHS )
  END IF
C Destroy the instance (deallocate internal data structures)
  id%JOB = -2
  CALL DMUMPS(id)
  CALL MPI_FINALIZE(IERR)
  STOP
END

```

Figure 4: Example program using MUMPS on an elemental DOUBLE PRECISION problem.

```

/* Example program using the C interface to the
 * double precision version of MUMPS, dmumps_c.
 * We solve the system  $A x = RHS$  with
 *  $A = \text{diag}(1 \ 2)$  and  $RHS = [1 \ 4]^T$ 
 * Solution is  $[1 \ 2]^T$  */
#include <stdio.h>
#include "mpi.h"
#include "dmumps_c.h"
#define JOB_INIT -1
#define JOB_END -2
#define USE_COMM_WORLD -987654
int main(int argc, char ** argv) {
    DMUMPS_STRUC_C id;
    int n = 2;
    int nz = 2;
    int irn[] = {1,2};
    int jcn[] = {1,2};
    double a[2];
    double rhs[2];

    int myid, ierr;
    ierr = MPI_Init(&argc, &argv);
    ierr = MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    /* Define A and rhs */
    rhs[0]=1.0;rhs[1]=4.0;
    a[0]=1.0;a[1]=2.0;

    /* Initialize a MUMPS instance. Use MPI_COMM_WORLD. */
    id.job=JOB_INIT; id.par=1; id.sym=0;id.comm_fortran=USE_COMM_WORLD;
    dmumps_c(&id);
    /* Define the problem on the host */
    if (myid == 0) {
        id.n = n; id.nz =nz; id.irn=irn; id.jcn=jcn;
        id.a = a; id.rhs = rhs;
    }
#define ICNTL(I) icntl[(I)-1] /* macro s.t. indices match documentation */
/* No outputs */
    id.ICNTL(1)=-1; id.ICNTL(2)=-1; id.ICNTL(3)=-1; id.ICNTL(4)=0;
/* Call the MUMPS package. */
    id.job=6;
    dmumps_c(&id);
    id.job=JOB_END; dmumps_c(&id); /* Terminate instance */
    if (myid == 0) {
        printf("Solution is : (%8.2f %8.2f)\n", rhs[0],rhs[1]);
    }
    return 0;
}

```

Figure 5: Example program using MUMPS from C on an assembled problem.

11 Notes on MUMPS distribution

This version of MUMPS is provided to you free of charge. It is public domain, based on public domain software developed during the Esprit IV European project PARASOL (1996-1999) by CERFACS, ENSEEIHT-IRIT and RAL. Since this first public domain version in 1999, the developments are supported by the following institutions: CERFACS, ENSEEIHT-IRIT, and INRIA Rhone-Alpes.

Main contributors are Patrick Amestoy, Iain Duff, Abdou Guermouche, Jacko Koster, Jean-Yves L'Excellent, and Stephane Pralet.

Up-to-date copies of the MUMPS package can be obtained from the Web pages <http://www.enseeiht.fr/apo/MUMPS/> or <http://graal.ens-lyon.fr/MUMPS>

THIS MATERIAL IS PROVIDED AS IS, WITH ABSOLUTELY NO WARRANTY EXPRESSED OR IMPLIED. ANY USE IS AT YOUR OWN RISK.

User documentation of any code that uses this software can include this complete notice. You can acknowledge (using references [1], [2], and [3] the contribution of this package in any scientific publication dependent upon the use of the package. You shall use reasonable endeavours to notify the authors of the package of this publication.

[1] P. R. Amestoy, I. S. Duff and J.-Y. L'Excellent (1998), Multifrontal parallel distributed symmetric and unsymmetric solvers, *Comput. Methods in Appl. Mech. Eng.*, 184, 501-520 (2000).

[2] P. R. Amestoy, I. S. Duff, J. Koster and J.-Y. L'Excellent, A fully asynchronous multifrontal solver using distributed dynamic scheduling, *SIAM Journal of Matrix Analysis and Applications*, Vol 23, No 1, pp 15-41 (2001).

[3] P. R. Amestoy, A. Guermouche, J.-Y. L'Excellent and S. Pralet (2005), Hybrid scheduling for the parallel solution of linear systems, *Parallel Computing* Vol 32, No 2, pp 136--156 (2006).

Accepted to *Parallel Computing*.

Other acknowledgements

Apart from the main contributors cited above, we are also grateful to Caroline Bousquet, Christophe Daniel, Vincent Espirat, Aurélia Fèvre, Chiara Puglisi, Grégoire Richard, Miroslav Tůma and Christophe Vömel who have been contributing to this work.

We are grateful to Jürgen Schulze for letting us distribute PORD developed at the University of Paderborn.

We also want to thank BRGM, EADS-CCR, INRIA-ScAIApplix, LARIA, Lawrence Berkeley National Laboratory, PARALLAB (Bergen) and Rutherford Appleton Laboratory for research discussions

that have certainly influenced this work.

Finally we want to thank the institutions that have provided access to their parallel machines: Centre Informatique National de l'Enseignement Supérieur (CINES), CERFACS, CICT (Toulouse), Fédération Lyonnaise de Calcul Haute-Performance, Institut du Développement et des Ressources en Informatique Scientifique (IDRIS), Lawrence Berkeley National Laboratory, Laboratoire de l'Informatique du Parallélisme, INRIA Rhones-Alpes, PARALLAB.

References

- [1] P. R. Amestoy. Recent progress in parallel multifrontal solvers for unsymmetric sparse matrices. In *Proceedings of the 15th World Congress on Scientific Computation, Modelling and Applied Mathematics, IMACS 97, Berlin*, 1997.
- [2] P. R. Amestoy, T. A. Davis, and I. S. Duff. An approximate minimum degree ordering algorithm. *SIAM Journal on Matrix Analysis and Applications*, 17:886–905, 1996.
- [3] P. R. Amestoy and I. S. Duff. Vectorization of a multiprocessor multifrontal code. *International Journal of Supercomputer Applications*, 3:41–59, 1989.
- [4] P. R. Amestoy, I. S. Duff, J. Koster, and J.-Y. L’Excellent. A fully asynchronous multifrontal solver using distributed dynamic scheduling. *SIAM Journal on Matrix Analysis and Applications*, 23(1):15–41, 2001.
- [5] P. R. Amestoy, I. S. Duff, and J.-Y. L’Excellent. Multifrontal solvers within the PARASOL environment. In B. Kågström, J. Dongarra, E. Elmroth, and J. Waśniewski, editors, *Applied Parallel Computing, PARA’98*, Lecture Notes in Computer Science, No. 1541, pages 7–11, Berlin, 1998. Springer-Verlag.
- [6] P. R. Amestoy, I. S. Duff, and J.-Y. L’Excellent. Parallélisation de la factorisation LU de matrices creuses non-symétriques pour des architectures à mémoire distribuée. *Calculateurs Parallèles Réseaux et Systèmes Répartis*, 10(5):509–520, 1998.
- [7] P. R. Amestoy, I. S. Duff, and J.-Y. L’Excellent. Multifrontal parallel distributed symmetric and unsymmetric solvers. *Comput. Methods Appl. Mech. Eng.*, 184:501–520, 2000.
- [8] P. R. Amestoy, A. Guermouche, J.-Y. L’Excellent, and S. Pralet. Hybrid scheduling for the parallel solution of linear systems. *Parallel Computing*, 32(2):136–156, 2006.
- [9] M. Arioli, J. Demmel, and I. S. Duff. Solving sparse linear systems with sparse backward error. *SIAM Journal on Matrix Analysis and Applications*, 10:165–190, 1989.
- [10] L. S. Blackford, J. Choi, A. Cleary, E. D’Azevedo, J. Demmel, I. Dhillon, J. Dongarra, S. Hammarling, G. Henry, A. Petitet, K. Stanley, D. Walker, and R. C. Whaley. *ScaLAPACK Users’ Guide*. SIAM Press, 1997.
- [11] A. R. Curtis and J. K. Reid. On the automatic scaling of matrices for Gaussian elimination. *J. Inst. Maths. Applics.*, 10:118–124, 1972.
- [12] J. J. Dongarra, J. Du Croz, I. S. Duff, and S. Hammarling. Algorithm 679. A set of Level 3 Basic Linear Algebra Subprograms. *ACM Transactions on Mathematical Software*, 16:1–17, 1990.
- [13] J. J. Dongarra, J. Du Croz, I. S. Duff, and S. Hammarling. Algorithm 679. A set of Level 3 Basic Linear Algebra Subprograms: model implementation and test programs. *ACM Transactions on Mathematical Software*, 16:18–28, 1990.
- [14] I. S. Duff and J. Koster. The design and use of algorithms for permuting large entries to the diagonal of sparse matrices. *SIAM Journal on Matrix Analysis and Applications*, 20(4):889–901, 1999.
- [15] I. S. Duff and J. Koster. On algorithms for permuting large entries to the diagonal of a sparse matrix. *SIAM Journal on Matrix Analysis and Applications*, 22(4):973–996, 2001.
- [16] I. S. Duff and J. K. Reid. The multifrontal solution of indefinite sparse symmetric linear systems. *ACM Transactions on Mathematical Software*, 9:302–325, 1983.
- [17] I. S. Duff and J. K. Reid. The multifrontal solution of unsymmetric sets of linear systems. *SIAM Journal on Scientific and Statistical Computing*, 5:633–641, 1984.
- [18] Iain S. Duff and Stéphane Pralet. Strategies for scaling and pivoting for sparse symmetric indefinite problems. *SIAM Journal on Matrix Analysis and Applications*, 27(2):313–340, 2005.
- [19] Aurélia Fèvre, Jean-Yves L’Excellent, and Stéphane Pralet. Scilab and MATLAB interfaces to MUMPS. Technical Report RR-5816, INRIA, January 2006. Also appeared as ENSEEIHT-IRIT report TR/TLSE/06/01 and LIP report RR2006-06.
- [20] A. Guermouche and J.-Y. L’Excellent. Constructing memory-minimizing schedules for multifrontal methods. *ACM Transactions on Mathematical Software*, 32(1):17–32, 2006.

- [21] A. Guermouche, J.-Y. L'Excellent, and G. Utard. Impact of reordering on the memory of a multifrontal solver. *Parallel Computing*, 29(9):1191–1218, 2003.
- [22] G. Karypis and V. Kumar. *METIS – A Software Package for Partitioning Unstructured Graphs, Partitioning Meshes, and Computing Fill-Reducing Orderings of Sparse Matrices – Version 4.0*. University of Minnesota, September 1998.
- [23] J. Schulze. Towards a tighter coupling of bottom-up and top-down sparse matrix ordering methods. *BIT*, 41(4):800–841, 2001.
- [24] M. Snir, S. W. Otto, S. Huss-Lederman, D. W. Walker, and J. Dongarra. *MPI: The Complete Reference*. The MIT Press, Cambridge, Massachusetts, 1996.